

5118007-02 Computer Architecture

Ch. 3 Arithmetic for Computers

9 Apr 2024

Shin Hong

Integer Arithmetic

- addition and subtraction
 - dealing with overflow
- multiplication
- division

Addition

- Digits are added bit-by-bit from right to left, considering carries passed from the right

$$\begin{array}{r} 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0111_{\text{two}} = 7_{\text{ten}} \\ + \quad 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0110_{\text{two}} = 6_{\text{ten}} \\ \hline = \quad 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 1101_{\text{two}} = 13_{\text{ten}} \end{array}$$

Subtraction

- A subtraction is basically the same as a combination of negation and addition

$$\begin{array}{r} 0000 \quad 0111_{\text{two}} = 7_{\text{ten}} \\ + \quad 1111 \quad 1010_{\text{two}} = -6_{\text{ten}} \\ \hline = \quad 0000 \quad 0001_{\text{two}} = 1_{\text{ten}} \end{array}$$

Overflow (1/2)

- An overflow occurs when the resulting value cannot be represented within given bits
 - only when adding two numbers of the same sign, and not possible when two numbers of different signs are added
- Overflow for signed numbers
 - adding two positive numbers yields a negative number
 - subtracting two negative numbers yields a positive number
 - the occurrence of overflow with add, addi, sub instructions causes an exception

Overflow (2/2)

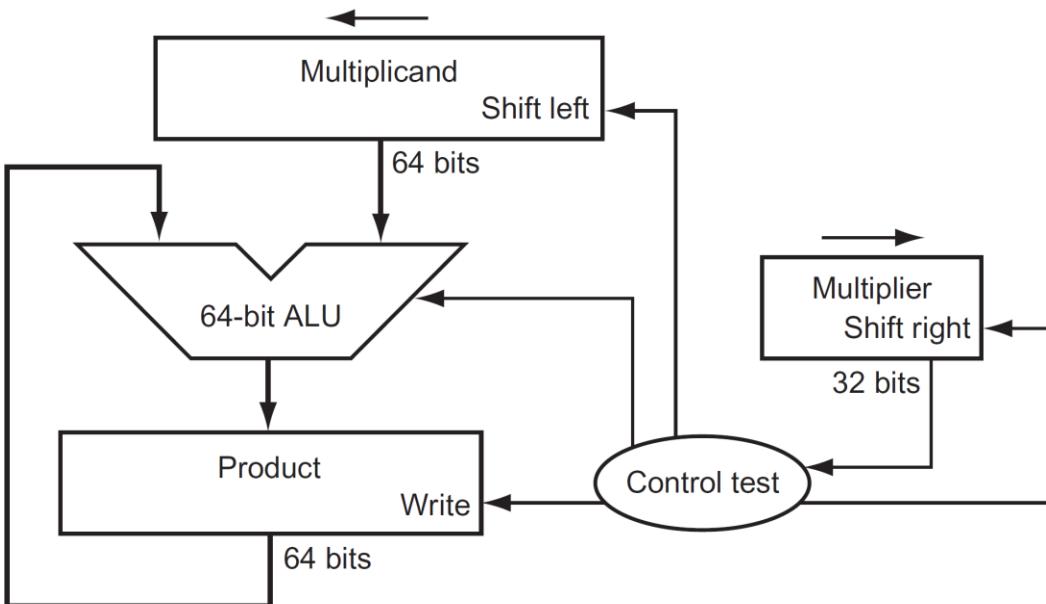
- Overflow for signed numbers
 - addition of two numbers yields a smaller number
 - subtraction of two numbers yields a greater number
 - overflow with addu, addiu, subu do not cause exception, thus C compiler uses these instructions for arithmetics.

Multiplication

$$\begin{array}{r} \text{Multiplicand} & 1000_{\text{ten}} \\ \text{Multiplier} & \times \quad \underline{1001_{\text{ten}}} \\ & 1000 \\ & 0000 \\ & 0000 \\ & 1000 \\ \hline \text{Product} & 1001000_{\text{ten}} \end{array}$$

- Basically, multiplication can be implemented as a sequential operations of right-shifts, left-shifts and additions
- Multiplication of a n -bits multiplicand and a m -bits multiplier produces a $n+m$ product
 - overflow may occur

Sequential Multiplication Hardware




```
foreach bit of multiplier begin
    if the LSB of multiplier is 1 then
        add multiplicand and product;
    end-if
    shift-left the multiplicand ;
    shift-right a multiplier ;
end-for
```

Example

- Multiply $0010_{(2)}$ and $0011_{(2)}$

Iteration	Step	Multiplier	Multiplicand	Product
0	Initial values	001①	0000 0010	0000 0000
1	1a: 1 \Rightarrow Prod = Prod + Mcand	0011	0000 0010	0000 0010
	2: Shift left Multiplicand	0011	0000 0100	0000 0010
	3: Shift right Multiplier	000①	0000 0100	0000 0010
2	1a: 1 \Rightarrow Prod = Prod + Mcand	0001	0000 0100	0000 0110
	2: Shift left Multiplicand	0001	0000 1000	0000 0110
	3: Shift right Multiplier	000①	0000 1000	0000 0110
3	1: 0 \Rightarrow No operation	0000	0000 1000	0000 0110
	2: Shift left Multiplicand	0000	0001 0000	0000 0110
	3: Shift right Multiplier	000①	0001 0000	0000 0110
4	1: 0 \Rightarrow No operation	0000	0001 0000	0000 0110
	2: Shift left Multiplicand	0000	0010 0000	0000 0110
	3: Shift right Multiplier	000①	0010 0000	0000 0110

Refined Version



Signed Multiplication

- Naive approach
 - store the signs of multiplicand and multiplier respectively
 - convert them to positive numbers
 - perform multiplication
 - apply the proper sign to the product
- Interestingly, the aforementioned multiplication algorithm works sound for two's complements
 - https://pages.cs.wisc.edu/~markhill/cs354/Fall2008/beyond354/int_mult.html

No Thinking Method

- sign-extend to both multiplicand and multiplier before multiplication

$$\begin{array}{r} 1111 \ 1111 \\ \times \ 1111 \ 1001 \\ \hline 11111111 \\ 00000000 \\ 00000000 \\ 11111111 \\ 11111111 \\ 11111111 \\ + \ 11111111 \\ \hline 1 \ 00000000111 \end{array}$$

$$\begin{array}{r} -1 \\ \times \ -7 \\ \hline 7 \end{array}$$

WRONG !
0011 (3)
 $\times 1011 (-5)$

0011
0011
0000
+ 0011


0100001
not -15 in any representation!

Sign extended:
0000 0011 (3)
 $\times 1111 \ 1011 (-5)$

00000011
00000011
00000000
00000011
00000011
00000011
00000011
+ 00000011

1011110001

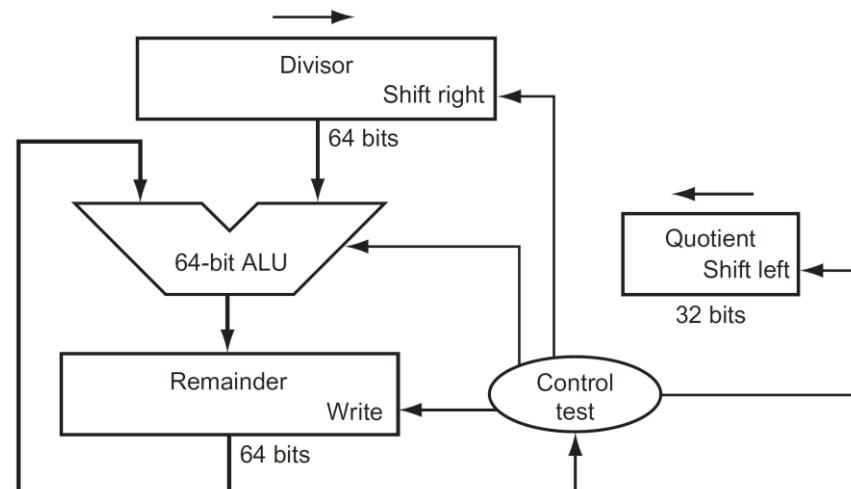
Fast Multiplication with Large Circuits



Division

- Division can be implemented as a series of shift-right, shift-left and subtraction


$$\begin{array}{r} \text{Divisor } 1000_{\text{ten}} \quad \text{Quotient} \\ \overline{)1001010_{\text{ten}}} \quad \text{Dividend} \\ -1000 \\ \hline 10 \\ 101 \\ 1010 \\ -1000 \\ \hline 10_{\text{ten}} \quad \text{Remainder} \end{array}$$



Example: $0111_{(2)}$ / $0010_{(2)}$

Iteration	Step	Quotient	Divisor	Reminder
0	Initial values	0000	0010 0000	0000 0111
1	1: Rem = Rem – Div	0000	0010 0000	①110 0111
	2b: Rem < 0 \Rightarrow +Div, sll Q, Q0 = 0	0000	0010 0000	0000 0111
	3: Shift Div right	0000	0001 0000	0000 0111
2	1: Rem = Rem – Div	0000	0001 0000	①111 0111
	2b: Rem < 0 \Rightarrow +Div, sll Q, Q0 = 0	0000	0001 0000	0000 0111
	3: Shift Div right	0000	0000 1000	0000 0111
3	1: Rem = Rem – Div	0000	0000 1000	①111 1111
	2b: Rem < 0 \Rightarrow +Div, sll Q, Q0 = 0	0000	0000 1000	0000 0111
	3: Shift Div right	0000	0000 0100	0000 0111
4	1: Rem = Rem – Div	0000	0000 0100	②000 0011
	2a: Rem \geq 0 \Rightarrow sll Q, Q0 = 1	0001	0000 0100	0000 0011
	3: Shift Div right	0001	0000 0010	0000 0011
5	1: Rem = Rem – Div	0001	0000 0010	②000 0001
	2a: Rem \geq 0 \Rightarrow sll Q, Q0 = 1	0011	0000 0010	0000 0001
	3: Shift Div right	0011	0000 0001	0000 0001

Faster Implementation



Floating Point (1/2)

- A real number can be represented as a combination of an integer and a fraction
- Scientific notation for decimal numbers
 - a single digit to the left of the decimal point and a power of ten
 - ex.
 - 1.02×10^{-9}
 - 3.15576×10^9

Floating Point (2/2)


- Scientific notation of binary number
 - $1.abcd \times 2^n$
 - $1 \times 2^n + a \times 2^{n-1} + b \times 2^{n-2} + c \times 2^{n-3} + d \times 2^{n-4}$
- Convert a decimal real number to binary
 - Examples
 - 4.5
 - 25.25
 - 32.45

Floating Number Representation (1/2)

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
s	exponent																														
1 bit	8 bits																														23 bits

- $(-1)^S \times F \times 2^E$
 - S: sign
 - F: fraction
 - E: exponent
- Single-precision: 8 bits for E and 23 bits for F
 - the smallest possible non-zero positive: 2.0×10^{-38}
 - the greatest possible positive: 2.0×10^{38}

Floating Number Representation (2/2)



- Double precision: 11 bits for E and 52 bits for F
 - the smallest possible non-zero positive: 2.0×10^{-308}
 - the greatest possible positive: 2.0×10^{308}

IEEE 754 Standard of Floating Number Encoding

Single precision		Double precision		Object represented
Exponent	Fraction	Exponent	Fraction	
0	0	0	0	0
0	Nonzero	0	Nonzero	\pm denormalized number
1–254	Anything	1–2046	Anything	\pm floating-point number
255	0	2047	0	\pm infinity
255	Nonzero	2047	Nonzero	NaN (Not a Number)

- Special symbols
 - Infinity
 - Not a Number (NaN)

Biased Notation

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0		
0	1	1	1	1	1	1	1	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0		
0	0	0	0	0	0	0	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

- Exponent does not use the two's complement notation, but a biased notation to make real number comparison similar to that of integer
- If an actual exponent is X with bias B, then represent it as X + B
 - Ex. for single precision with bias 127, if an exponent is -1, it is represented as $0111\ 1110_{(\text{two})} = 126 = (-1 + 127)$

Example

- Represent -0.75_{ten} in IEEE 754 single and double precision

The number -0.75_{ten} is also

$$-3/4_{\text{ten}} \text{ or } -3/2^2_{\text{ten}}$$

It is also represented by the binary fraction

$$-11_{\text{two}}/2^2_{\text{ten}} \text{ or } -0.11_{\text{two}}$$

In scientific notation, the value is

$$-0.11_{\text{two}} \times 2^0$$

and in normalized scientific notation, it is

$$-1.1_{\text{two}} \times 2^{-1}$$

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
1	0	1	1	1	1	1	1	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

1 bit 8 bits 23 bits

Example

- Represent $-0.75_{(\text{ten})}$ in IEEE 754 single and double precision

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
1	0	1	1	1	1	1	1	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

1 bit 8 bits 23 bits

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
1	0	1	1	1	1	1	1	1	1	1	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

1 bit 11 bits 20 bits

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

32 bits

Example

- What decimal number it is?

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
1	1	0	0	0	0	0	0	1	0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	.	.	.

The sign bit is 1, the exponent field contains 129, and the fraction field contains $1 \times 2^{-2} = 1/4$, or 0.25. Using the basic equation,

$$\begin{aligned}(-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent}-\text{Bias})} &= (-1)^1 \times (1 + 0.25) \times 2^{(129-127)} \\&= -1 \times 1.25 \times 2^2 \\&= -1.25 \times 4 \\&= -5.0\end{aligned}$$

Floating-point Addition

- Example

$$9.999_{\text{ten}} \times 10^1 + 1.610_{\text{ten}} \times 10^{-1}.$$

Binary Floating-Point Addition

Try adding the numbers 0.5_{ten} and -0.4375_{ten} in binary using the algorithm in Figure 3.14.

Let's first look at the binary version of the two numbers in normalized scientific notation, assuming that we keep 4 bits of precision:

$$\begin{array}{lll} 0.5_{\text{ten}} & = 1/2_{\text{ten}} & = 1/2^1_{\text{ten}} \\ & = 0.1_{\text{two}} & = 0.1_{\text{two}} \times 2^0 & = 1.000_{\text{two}} \times 2^{-1} \\ -0.4375_{\text{ten}} & = -7/16_{\text{ten}} & = -7/2^4_{\text{ten}} \\ & = -0.0111_{\text{two}} & = -0.0111_{\text{two}} \times 2^0 & = -1.110_{\text{two}} \times 2^{-2} \end{array}$$

Now we follow the algorithm:

Step 1. The significand of the number with the lesser exponent ($-1.11_{\text{two}} \times 2^{-2}$) is shifted right until its exponent matches the larger number:

$$-1.110_{\text{two}} \times 2^{-2} = -0.111_{\text{two}} \times 2^{-1}$$

Step 2. Add the significands:

$$1.000_{\text{two}} \times 2^{-1} + (-0.111_{\text{two}} \times 2^{-1}) = 0.001_{\text{two}} \times 2^{-1}$$

Step 3. Normalize the sum, checking for overflow or underflow:

$$\begin{aligned}0.001_{\text{two}} \times 2^{-1} &= 0.010_{\text{two}} \times 2^{-2} = 0.100_{\text{two}} \times 2^{-3} \\&= 1.000_{\text{two}} \times 2^{-4}\end{aligned}$$

Since $127 \geq +4 \geq -126$, there is no overflow or underflow. (The biased exponent would be $-4 + 127$, or 123, which is between 1 and 254, the smallest and largest unreserved biased exponents.)

Step 4. Round the sum:

$$1.000_{\text{two}} \times 2^{-4}$$

The sum already fits exactly in 4 bits, so there is no change to the bits due to rounding.

This sum is then

$$\begin{aligned}1.000_{\text{two}} \times 2^{-4} &= 0.0001000_{\text{two}} = 0.0001_{\text{two}} \\&= 1/2^4_{\text{ten}} = 1/16_{\text{ten}} = 0.0625_{\text{ten}}\end{aligned}$$

This sum is what we would expect from adding 0.5_{ten} to -0.4375_{ten} .