



Lecture 11: Searching algorithm (Part. 2)

Algorithm

Jeong-Hun Kim

Table of Contents

❖ Part 1

- Red-black tree

❖ Part 2

- B-tree

❖ Part 3

- Multidimensional search tree
 - KD-tree
 - KDB-tree
 - R-tree
 - Grid file

Part 1

RED-BLACK TREE

Red-Black Tree

- ❖ Limitations of the binary search tree
 - Average time complexities for storage and search is $\Theta(\log n)$
 - In the worst case, the tree becomes unbalanced
 - When unbalanced, the time complexity is close to $\Theta(n)$
- ❖ Balanced tree
 - This always remains balanced
 - e.g., Red-Black tree, B-tree, AVL tree, etc.


Red-Black Tree

❖ Properties of a Red-Black tree


- Every node has either a **Red** or Black color
 - ① Every node is either **Red** or Black
 - ② Root node is Black
 - ③ All leaf nodes are Black
 - ④ If a node is **Red**, its children must be Black (No double **Red**)
 - ⑤ The number of Black nodes encountered on any path from the root node to a leaf node is the same
- Here, the leaf node of Red-Black tree is different from a traditional leaf node
 - Assume all NIL pointers point to a leaf node called NIL

Red-Black Tree


- ❖ Example of Red-Black tree



Binary search tree



Red-Black tree





Implemented Red-Black tree

- ① Root node is Black
- ② All leaf nodes are Black
- ③ If a node is Red, its children must be Black
- ④ The number of Black nodes encountered on any path from the root node to a leaf node is the same


Red-Black Tree

- ❖ Insertion in a Red-Black tree
 - Fundamentally the same as a binary search tree
 - Mark the inserted node (node x) as **Red** after insertion
 - If the color of node x 's parent node p is
 - » Black: no problem
 - » Red: property ③ of the Red-Black tree is not satisfied



Red-Black Tree

- ❖ Insertion in a Red-Black tree (cont'd)
 - The node p^2 and the sibling node of x must be Black
 - Divided into two cases depending on the color of node s
 - Case 1: node s is Red
 - Case 2: node s is Black



Red-Black Tree

❖ Insertion in a Red-Black tree (cont'd)



- Case 1: node s is Red

- Change node p and its sibling node s to Black
 - However, the same issue may occur at node p^2
 - Solve recursively



: nodes with changed colors

Case 1





Red-Black Tree

❖ Insertion in a Red-Black tree (cont'd)

- Case 2: node s is Black

- Case 2-1: node s is Black, and node x is the **right child** of parent node p
 - Case 2-2: node s is Black, and node x is the **left child** of parent node p




Red-Black Tree


❖ Insertion in a Red-Black tree (cont'd)

- Case 2-1: node s is Black, and x is the **right child** of parent node p
 - Rotate left around node p
 - Violates property ③, transition to Case 2-2
 - New parent is Black, and the children are **Red**

Case 2-1



Case 2-2




Red-Black Tree

❖ Insertion in a Red-Black tree (cont'd)

- Case 2-2: node s is Black, and x is the **left child** of parent node p
 - Rotate right around node p^2 , and swap the colors of p and p^2


● : nodes with changed colors



Red-Black Tree

❖ Deletion in a Red-Black tree


- Deletion is straightforward if property ④ is satisfied
 - If the target is Red, there is no problem
 - Even if the target is Black, there is no problem if the only child is Red




Red-Black Tree

❖ Deletion in a Red-Black tree (cont'd)

- Deletion becomes complex if property ④ is violated
 - '-1' next to node x indicates that the number of Black nodes on the path from the root to the leaf through x is one less than required
 - Handled by dividing into a total of five cases



Problem occurs after deleting node m
(Property ④ is violated)



Handling method varies depending on
 x 's surrounding conditions








Red-Black Tree

❖ Deletion in a Red-Black tree (cont'd)

- Case 1: if node p is **Red** (node s must be Black) and depending on the <colors of l and r >
 - Case 1-1: <Black, Black>
 - Case 1-2: <*, Red>
 - Case 1-3: <Red, Black>
- Case 2: if node p is **Black** and depending on the <colors of s , l , and r >
 - Case 2-1: <Black, Black, Black>
 - Case 2-2: <Black, *, Red>
 - Case 2-3: <Black, Red, Black>
 - Case 2-4: <Red, Black, Black>
 - If s is Red, then l and r must be Black

Red-Black Tree

❖ Deletion in a Red-Black tree (cont'd)








Depending on the color of s

Depending on the color of p

Red-Black Tree

- ❖ Deletion in a Red-Black tree (cont'd)




Finally divided into five cases

Red-Black Tree

❖ Deletion in a Red-Black tree (cont'd)

- Case 1-1: swap the colors of nodes p and s
 - Number of Black nodes on the path to x increases
 - Number of Black nodes on the path through s remains unchanged




Red-Black Tree

❖ Deletion in a Red-Black tree (cont'd)

- Case *-2:

- ① Perform a left rotation around p
- ② Swap the colors of p and s
- ③ Change the color of r from Red to Black
 - Number of Black nodes on the path to x increases
 - Number of Black nodes on the path through s remains unchanged




Red-Black Tree

❖ Deletion in a Red-Black tree (cont'd)

- Case *-3:

- ① Perform a right rotation around s
- ② Swap the colors of l and r
- ③ Transition to Case *-2
 - Number of Black nodes on the path to x increases
 - Number of Black nodes on the path through s remains unchanged




Red-Black Tree

❖ Deletion in a Red-Black tree (cont'd)

- Case 2-1:

- ① Change the color of s from Black to Red
 - Path through s becomes deficient in Black nodes
 - Entire path through p becomes deficient in Black nodes
- ② Set p as the target node and solve recursively




Red-Black Tree

❖ Deletion in a Red-Black tree (cont'd)

- Case 2-4:

- ① Perform a left rotation around p
- ② Swap the colors of p and s
 - Paths through l and r have no issues
 - However, the color of x 's parent node changes from Black to Red
- ③ Consider the color combinations and transition to one of Case 1-1, 1-2, or 1-3



Red-Black Tree

❖ Analysis of the Red-Black tree

- Possible maximum depth of a Red-Black tree with a total of n keys:
 - $O(\log n)$
 - If the number of keys is n , then there are n internal nodes
 - The depth of the tree is $\lfloor \log n \rfloor + 1$
 - No matter how well constructed, the Black nodes on any path from the root node to an arbitrary leaf cannot exceed $\lfloor \log n \rfloor + 1$
 - Since Red nodes cannot exist consecutively, there are fewer Red nodes than Black nodes
 - The total length of any leaf cannot exceed $2(\lfloor \log n \rfloor + 1)$

Part 2

B-TREE

B-Tree

❖ Background

- Utilize resources used for disk access
 - Accessing a disk takes significantly more time compared to accessing memory
 - Accessing data on a disk requires block-level access for reading and writing
 - Even if you want to write just one byte, the entire block must be read first
 - Blocks are usually 8KB or 16KB
 - At the software level, a block is referred to as a page
 - Reading one block requires 200,000 clock cycles (based on a 2GHz CPU)
 - » Processing a machine instruction takes less than 2 clock cycles

B-Tree

❖ Background

- What if the search tree is too large?
 - Search tree cannot be loaded into memory for use
 - The search tree is placed on the disk for processing
 - Number of disk accesses has a greater impact on performance than CPU clock cycles
 - Increasing the branching factor of the search tree can reduce its expected depth
 - » What if handling around one billion keys?
 - » In the case of a binary tree, the depth is 30
 - » If there are 256 branches, the depth is 5
 - » If data access requests are frequent, the number of disk accesses is reduced to 1/6
 - » A branch refers to the number of children a parent can have

B-Tree


❖ Background

- External search tree
 - Search tree on the disk
- B-tree
 - A search tree with more than two branches
 - Suitable for use in a disk environment
 - Maintains balance to reduce the number of disk accesses

B-Tree

❖ B-tree

- Properties
 - All nodes except the root have $\lfloor k/2 \rfloor \sim k - 1$ keys
 - All leaf nodes have the same depth
- Maximize the number of branches while ensuring each node has at least half of its maximum allowable keys




B-Tree

- ❖ B-Tree of Order k has the following properties
 - ① All the leaf nodes must be at the same level
 - ② All nodes except root must have at least $\lfloor k/2 \rfloor$ keys and maximum of $k - 1$ keys
 - ③ All internal nodes except root must have at least $\lfloor k/2 \rfloor$ children
 - ④ If the root node is not leaf node, then it must have at least 2 children
 - ⑤ An internal node with $n - 1$ keys must have n number of children
 - ⑥ All keys within a node must be in ascending (descending) order

B-Tree

- ❖ B-Tree of Order k has the following properties
 - Property ①: all the leaf nodes must be at the same level

B-Tree of Order 4




B-Tree

❖ B-Tree of Order k has the following properties

- Property ②: all nodes except root must have at least $\lfloor k/2 \rfloor$ keys and maximum of $k - 1$ keys

B-Tree of Order 4




B-Tree

❖ B-Tree of Order k has the following properties

- Property ③: all internal nodes except root must have at least $\lfloor k/2 \rfloor$ children

B-Tree of Order 4




B-Tree

❖ B-Tree of Order k has the following properties

- Property ④: if the root node is not leaf, then it must have at least 2 children


B-Tree of Order 4



B-Tree

- ❖ B-Tree of Order k has the following properties
 - Property ⑤: an internal node with $n - 1$ keys must have n number of children


B-Tree of Order 4



B-Tree

- ❖ B-Tree of Order k has the following properties
 - Property ⑥: all keys within a node must be in ascending (descending) order

B-Tree of Order 4



B-Tree

- ❖ In a B-tree, new element must be added only at leaf node
 - ① If tree is **empty**, then create a new node with new key and insert into the tree as a root node
 - ② If tree is **not empty**, then find a leaf node to which the new key can be added using binary search tree logic
 - ③ If the leaf node has an empty space, then add the new key to the leaf node by maintaining **order of keys** within the node
 - ④ If the leaf node is already full, then split the node by sending middle value to its parent node. Repeat that same until sending value is fixed into a node
 - ⑤ If the splitting is occurring to the root node, then the middle value becomes **new root node** for the tree and the height of the tree is increased by one

B-Tree

❖ Example

- Construct a B-tree of Order 3 ($k=3$) by inserting numbers from 1 to 10
- Insert '1'
 - Rule ①:
 - If tree is empty, then create a new node with new key and insert into the tree as a root node



B-Tree

❖ Example

- Insert '2'
 - Rule ②:
 - If tree is not empty, then find a leaf node to which the new key can be added using binary search tree logic
 - Rule ③:
 - If the leaf node has an empty space, then add the new key to the leaf node by maintaining order of keys within the node



B-Tree

❖ Example

- Insert '3'

- Rule ④:

- If the leaf node is already full, then split the node by sending middle value to its parent node

- Rule ⑤:


- If the splitting is occurring to the root node, then the middle value becomes new root node for the tree and the height of the tree is increased by one




B-Tree

❖ Example

- Insert '4'




- Insert '5'




B-Tree

❖ Example

- Insert '6'




- Insert '7'




B-Tree

❖ Example

- Insert '8'




- Insert '9'




B-Tree

- ❖ Example
 - Insert '10'



B-Tree

- ❖ How to search a key using B-tree
 - B-tree searches for the key top-down from the root node to the leaf
 - Search '7'



B-Tree

❖ Deletion in a B-tree

- Process:

- ① Find the node with key x
- ② If it is not a leaf node, find the leaf node r with key y , where y is the element immediately following x
- ③ Delete x from the leaf node
- ④ Resolve any underflow in the node after deleting x

B-Tree

❖ Deletion in a B-tree

- Pseudo-code


```
Deletion(t, x, v){  
    if (v is not leaf) then{  
        find a leaf node with a key y  
        swap x and y  
    }  
    delete x in the leaf node r  
    if (underflow occurs in r) then{  
        clearUnderflow(r)  
    }  
}
```

```
clearUnderflow(r){  
    if (r's sibling nodes has extra keys) then  
        send the extra keys to r  
    else {  
        merge r with its sibling nodes  
        if (underflow occurs in the parent p)  
            then  
                clearUnderflow(p)  
    }  
}
```


B-Tree

❖ Deletion in a B-tree

- Example ($k = 6$)



Delete '7'




Delete '4'

B-Tree

❖ Deletion in a B-tree


- Example ($k = 5$)



B-Tree

❖ Deletion in a B-tree

- Example ($k = 5$)



B-Tree

❖ Analysis of the B-tree

- If a binary search tree is well balanced, its height approaches $\log_2 n$
- If a d-branches tree is well balanced, its height approaches $\log_d n$
- In a B-tree, if $k = d$, every internal node except the root must have at least $\left\lfloor \frac{d}{2} \right\rfloor$ children
 - In worst case, the height of B-tree is $\log_{\left\lfloor \frac{d}{2} \right\rfloor} n$

B-Tree

❖ Analysis of the B-tree (cont'd)

- Searching time: $O(\log n)$
- Insertion time: $O(\log n)$
- Deletion time: $O(\log n)$

Part 3

MULTIDIMENSIONAL SEARCH TREE

Multidimensional Search Tree

❖ What is Multidimensional Search Tree?

- Single dimensional search tree
 - Record key consists of a single field
- Multidimensional search tree
 - Record key consists of multiple fields
 - Indexing for **spatial data**
- Representative trees
 - KD-tree
 - KDB-tree
 - R-tree
 - Grid file


Multidimensional Search Tree

❖ KD-tree

- An extension of the binary search tree
- Fields composing the keys are used alternately for searching at each level
 - Root node branches using only the first field
 - The next level branches using only the second field
 - (...)
 - The i -th level branches using only the $i \bmod k$ field

Multidimensional Search Tree


- ❖ Example of KD-tree



Multidimensional Search Tree

❖ Example of KD-tree



- Two-dimensional KD-tree represented by (x, y)
 - Input sequence: A(50, 50) → B(10, 70) → C(80, 85) → D(25, 20) → E(40, 85) → F(70, 85) → G(10, 60)



Multidimensional Search Tree

❖ KD-tree and multidimensional space

- KD-tree is a process of partitioning space
- KD-tree performs data searches by gradually narrowing down the space



Multidimensional Search Tree

- ❖ Searching in a KD-tree
 - ① Input an arbitrary key
 - ② Use each field in sequence for the search
- ❖ Insertion in a KD-tree
 - ① Input the new key to insert
 - ② Use each field in sequence for the search
 - ③ Traverse the tree, and upon reaching a leaf node, insert on the left or right


Searching and inserting are simple extensions of a binary search tree

Multidimensional Search Tree

- ❖ Deletion in a KD-tree
 - Binary search tree
 - Handle by categorizing into three cases based on the number of children
 - KD-tree
 - Case 1: no children
 - Remove only the target node r
 - Case 2: one children
 - Directly connect the parent of the target node r to r 's child node
 - However, the properties of the KD-tree are not preserved
 - Therefore, delete it using the same method as when there are two children
 - Case 3: two children
 - Move the node with the smallest branching field value in the right subtree to the position of node r
 - If there is only a left child, move the node with the largest branching field value in the left subtree to the position of node r

Multidimensional Search Tree


- ❖ Deletion in a KD-tree
 - Example: delete node A



Multidimensional Search Tree

- ❖ Deletion in a KD-tree
 - Example: delete node A


(b)



Multidimensional Search Tree

- ❖ Deletion in a KD-tree
 - Example: delete node A


(c)



Multidimensional Search Tree

- ❖ Deletion in a KD-tree
 - Example: delete node A


(d)



Multidimensional Search Tree

- ❖ Deletion in a KD-tree
 - Example: delete node A


(e)



Multidimensional Search Tree


- ❖ Deletion in a KD-tree
 - Example: delete node A

(f)



Multidimensional Search Tree

- ❖ Deletion in a KD-tree
 - The parts of the KD-tree that do not need to be traversed when searching for the smallest value x



Multidimensional Search Tree

- ❖ Example uses of a KD-tree
 - Multidimensional k-nearest neighbor search
 - <https://www.quora.com/What-is-a-kd-tree-and-what-is-it-used-for>

1. Nearest Neighbor Search

Let's say you intend to build a **Social Cop** in your smartphone. Social Cop helps people report crimes to the nearest police station in real-time.

So what seems to be a problem here ?

Yes, you guessed it right. We need to search for the police station nearest to the crime location before attempting to report anything.


*How could we do it **quickly** ?*

Seems k-d trees can help you find the nearest neighbor to a point on a two dimensional map of your city. All you have to do is construct a 2 dimensional k-d tree from the locations of all the police stations in your city, and then query the k-d tree to find the nearest police station to any given location in the city.

Okay, I get what they can do. But how do they do it ?

If you already know how [binary search trees](#) work, understanding how k-d trees work would be nothing new. k-d trees help in partitioning space just as binary search trees help in partitioning the [real line](#). k-d trees recursively partition a region of space, creating a binary space partition at each level of the tree.

This is what a 3 dimensional region of space partitioned by a 3 dimensional k-d tree looks like [1] :



Multidimensional Search Tree


- ❖ Example uses of a KD-tree
 - Database query
 - <https://www.quora.com/What-is-a-kd-tree-and-what-is-it-used-for>

2. Database queries involving a multidimensional search key

A query asking for all the employees in the age-group of (40, 50) and earning a salary in the range of (15000, 20000) per month can be transformed into a geometrical problem where the age is plotted along the x-axis and the salary is plotted along the y-axis [4]

[4] The x-axis denotes the age of the employee in *years*, and the y-axis denotes the monthly salary in *thousand rupees*.

A 2 dimensional k-d tree on the composite index of (*age, salary*) could help you efficiently search for all the employees that fall in the rectangular region of space defined by the query described above.




Multidimensional Search Tree

❖ R-tree

- Extend B-tree to handle multidimensional searches
- Balanced tree
- Nodes in an R-tree represent the minimum bounding space that contains all relevant keys

이름	Key1	Key2
A	8	100
B	4	10
C	6	35
D	1	10
E	6	40
F	5	45
G	7	85
H	3	20
I	10	70
J	2	30
K	8	50
L	4	50



Multidimensional Search Tree

- ❖ Nodes in R-tree

- Region page
 - Composed of multiple (region, page number) pairs
 - All internal nodes are region pages
 - Point page
 - All leaf nodes are point pages


- ❖ Upper and lower bounds on the number of regions in an R-tree

- All internal nodes except the root have $\left\lfloor \frac{k}{2} \right\rfloor \sim k$ regions
 - All leaf nodes have the same depth
 - All records are referenced only in the leaf nodes

Multidimensional Search Tree

❖ Example of R-tree

- If $k=5$, all nodes except the root have 2 to 5 elements
- Leaf nodes contain the actual (key, page number) pairs
- Can be included in overlapping regions





이름	Key1	Key2
A	8	100
B	4	10
C	6	35
D	1	10
E	6	40
F	5	45
G	7	85
H	3	20
I	10	70
J	2	30
K	8	50
L	4	50

Multidimensional Search Tree

❖ Example of R-tree



- Insert M(5, 20)
- Insert N(4, 30)



Multidimensional Search Tree

❖ Example of R-tree



- Redistribution for inserting N(4, 30)



Multidimensional Search Tree

❖ Example of R-tree



- Insert O(3, 30) and P(2, 20)
- Insert Q(3, 10)



Multidimensional Search Tree


❖ Example of R-tree

- Redistribution for inserting Q(3, 10)



Multidimensional Search Tree

- ❖ Insertion and deletion in an R-tree
 - Similar to a B-tree
- ❖ In an R-tree, node regions can overlap
 - Allows flexibility in node adjustments
 - However, the path for record search may not be unique
 - This is improved in the R*-tree
- ❖ R-tree can also be used for storing and searching geometric shapes / data




Multidimensional Search Tree

- ❖ Grid file
 - Divides the space into mutually exclusive grid regions
 - Stores only the records belonging to that region
 - Allows for fast storage and retrieval of arbitrary records



Multidimensional Search Tree

- ❖ Example of Grid file




Multidimensional Search Tree

- ❖ Example of Grid file (cont'd)





Multidimensional Search Tree

- ❖ Example of Grid file (cont'd)





Multidimensional Search Tree

❖ Example of Grid file (cont'd)




Multidimensional Search Tree

❖ Example of Grid file (cont'd)



Multidimensional Search Tree

- ❖ Example of Grid file (cont'd)



Multidimensional Search Tree

❖ Searching in a Grid file

- If the grid array is not excessively large, it can be loaded into memory for use
 - However, this is not practically feasible
- If the grid array is not excessively large, it can be loaded into memory for use
 - Determine the page where the target record is stored in the grid array, and then go to that page for use

Summary

- ❖ Red-Black tree
 - Balanced tree
 - Properties:
 - Every node is either Red or Black
 - Root node is Black
 - All leaf nodes (NILs) are Black
 - No double Red
 - Same Black depth
- ❖ B-tree
 - Multiple branches
 - Useful as an external tree
 - Properties:
 - All nodes except the root have $\lfloor k/2 \rfloor \sim k - 1$ keys
 - All leaf nodes have the same depth
- ❖ Multidimensional search tree
 - KD-tree
 - KDB-tree
 - R-tree
 - Grid file

Questions?

SEE YOU NEXT TIME!