

# PROGRAM OPTIMIZATION

Jo, Heeseung

# Today

---

## Overview

### Generally Useful Optimizations

- Code motion/precomputation
- Strength reduction
- Sharing of common subexpressions
- Removing unnecessary procedure calls

### Optimization Blockers

- Procedure calls
- Memory aliasing

# Performance Realities

---

## Constant factors matter

- Easily see 10:1 performance range depending on how code is written
- Must optimize at multiple levels:
  - algorithm, data representations, procedures, and loops

## Must understand system to optimize performance

- How programs are compiled and executed
- How to measure program performance and identify bottlenecks
- How to improve performance without destroying code modularity and generality

# Optimizing Compilers

---

Provide efficient mapping of program to machine

- Register allocation
- Code selection and ordering (scheduling)
- Dead code elimination
- Eliminating minor inefficiencies

Up to programmer to select best overall algorithm

- Big-0 savings are (often) more important than constant factors
- But constant factors also matter

Have difficulty overcoming "optimization blockers"

- Potential memory aliasing
- Potential procedure side-effects

# Limitations of Optimizing Compilers

---

Operate under fundamental constraint

- Must not cause any change in program behavior

Most analysis is performed only **within procedures**

- Whole-program analysis is too expensive in most cases

Most analysis is **based only on static information**

- Compiler has difficulty anticipating run-time inputs

**When in doubt, the compiler must be conservative**

# Generally Useful Optimizations

## 1. Code motion

- Reduce frequency with which computation performed
  - If it will always produce same result
  - Especially moving code out of loop

```
void set_row(double *a, double *b,  
long i, long n)  
{  
    long j;  
    for (j = 0; j < n; j++)  
        a[n*i+j] = b[j];  
}
```



int  $ni = n * i$   $\Rightarrow$   
 $\Rightarrow$  곱하기를 ~~한번~~ 한 번  $\circ$  Loop ~~한 번~~ 한번.  
~~한 번~~ 한 번.  $\rightarrow$  한번 ~~한 번~~ 한번.

# Compiler-Generated Code Motion

```
void set_row(double *a, double *b,
    long i, long n)
{
    long j;
    for (j = 0; j < n; j++)
        a[n*i+j] = b[j];
}
```

```
long j;
long ni = n*i;
double *rowp = a+ni;
for (j = 0; j < n; j++)
    *rowp++ = b[j];
```

```
set_row:
    testq    %rcx, %rcx          # Test n
    jle      .L4                  # If 0, goto done
    movq    %rcx, %rax           # rax = n
    imulq   %rdx, %rax           # rax *= i
    leaq     (%rdi,%rax,8), %rdx # rowp = a + n*i*8
    movl    $0, %r8d              # j = 0
.L3:
    # loop:
    movq     (%rsi,%r8,8), %rax # t = b[j]
    movq     %rax, (%rdx)         # *rowp = t
    addq    $1, %r8                # j++
    addq    $8, %rdx               # rowp++
    cmpq    %r8, %rcx             # Compare n:j
    jg      .L3                  # If >, goto loop
.L4:
    rep ; ret
```

이제 여기서  
loop를 풀면 됩니다.  
⇒ code motion

# Generally Useful Optimizations

## 2. Reduction in strength

- Replace costly operation with simpler one
- Shift, add instead of multiply or divide
  - $16*x \rightarrow x \ll 4$
  - Utility machine dependent
  - Depends on cost of multiply or divide instruction
  - On Intel Nehalem, integer multiply requires 3 CPU cycles
- Recognize sequence of products

```
for (i = 0; i < n; i++)
    for (j = 0; j < n; j++)
        a[n*i + j] = b[j];
```

int ni = 0;

ni += n;



# Generally Useful Optimizations

## 3. Share common subexpressions

- Reuse portions of expressions
- Compilers often not very sophisticated in exploiting arithmetic properties

```
/* Sum neighbors of i,j */  
up = val[(i-1)*n + j];  
down = val[(i+1)*n + j];  
left = val[i*n      + j-1];  
right = val[i*n      + j+1];  
sum = up + down + left + right;
```

```
long inj = i*n + j;  
up = val[inj - n];  
down = val[inj + n];  
left = val[inj - 1];  
right = val[inj + 1];  
sum = up + down + left + right;
```

3 multiplications:  $i*n$ ,  $(i-1)*n$ ,  $(i+1)*n$

```
leaq  1(%rsi), %rax # i+1  
leaq -1(%rsi), %r8  # i-1  
imulq %rcx, %rsi    # i*n  
imulq %rcx, %rax    # (i+1)*n  
imulq %rcx, %r8    # (i-1)*n  
addq  %rdx, %rsi    # i*n+j  
addq  %rdx, %rax    # (i+1)*n+j  
addq  %rdx, %r8    # (i-1)*n+j
```

1 multiplication:  $i*n$

```
imulq  %rcx, %rsi    # i*n  
addq   %rdx, %rsi    # i*n+j  
movq   %rsi, %rax    # i*n+j  
subq   %rcx, %rax    # i*n+j-n  
leaq   (%rsi,%rcx), %rcx # i*n+j+n
```

# Optimization Blocker #1: Procedure Calls

Procedure to convert string to lower case

```
void lower(char *s)
{
    int i;
    for (i = 0; i < strlen(s); i++)
        if (s[i] >= 'A' && s[i] <= 'Z')
            s[i] -= ('A' - 'a');
}
```


What's wrong?

Nothing but wrong

# Lower Case Conversion Performance

Time quadruples when double string length

- Quadratic performance -> Why??




# Convert Loop To Goto Form

strlen executed every iteration

```
void lower(char *s)
{
    int i;
    for (i = 0; i < strlen(s); i++)
        if (s[i] >= 'A' && s[i] <= 'Z')
            s[i] -= ('A' - 'a');
}
```

// goto version of for statement

```
void lower(char *s)
{
    int i = 0;
    if (i >= strlen(s))
        goto done;
loop:   if (s[i] >= 'A' && s[i] <= 'Z')
        s[i] -= ('A' - 'a');
    i++;
    if (i < strlen(s))
        goto loop;
done:
}
```



# Calling strlen

## strlen() performance

- Only way to determine length of string is to scan its entire length, looking for null character

Overall performance, string of length N

- N calls to strlen
- Require times N
- Overall  $O(N^2)$  performance

```
size_t strlen(const char *s)
{
    size_t length = 0;
    while (*s != '\0') {
        s++;
        length++;
    }
    return length;
}
```

# Improving Performance

## Move call to strlen() outside of loop

- Since result does not change from one iteration to another
- Form of code motion

```
void lower(char *s)
{
    int i;
    for (i = 0; i < strlen(s); i++)
        if (s[i] >= 'A' && s[i] <= 'Z')
            s[i] -= ('A' - 'a');
}
```


```
void lower2(char *s)
{
    int i;
    int len = strlen(s);
    for (i = 0; i < len; i++)
        if (s[i] >= 'A' && s[i] <= 'Z')
            s[i] -= ('A' - 'a');
}
```

o(n)  $\approx$  2x.

strlen()의 결과는 헤더가 있을 수도 있고,  
~~~ 정기적 ~ ~ | 런타임에 찾을 수 있다.

# Lower Case Conversion Performance

Linear performance of lower2



# Optimization Blocker: Procedure Calls

Why couldn't compiler move `strlen()` out of inner loop?

- Procedure may have **side effects**
  - Alters **global state** each time called
- Function may not return same value for given arguments
  - Depends on other parts of global state
  - Procedure lower could interact with `strlen`

`strtok()`

Warning:

- Compiler treats procedure call as a black box
- Weak optimizations near them

Remedies:

- Use of inline functions
  - GCC does this with -O2
- Do your own code motion

```
int lencnt = 0;
size_t strlen(const char *s)
{
    size_t length = 0;
    while (*s != '\0') {
        s++; length++;
    }
    lencnt += length;
    return length;
}
```

# Memory Matters

Code updates  $b[i]$  on every iteration

- Why couldn't compiler optimize this?

```
/* Sum rows is of n X n matrix a  
and store in vector b */  
void sum_rows1(double *a, double *b, long n) {  
    long i, j;  
    for (i = 0; i < n; i++) {  
        b[i] = 0;  
        for (j = 0; j < n; j++)  
            b[i] += a[i*n + j];  
    }  
}
```

```
double A[9] =  
{ 0, 1, 2,  
 4, 8, 16,  
 32, 64, 128};  
double *B = {0, 0, 0};  
sum_rows1(A, B, 3);
```

Value of B:

init: [0, 0, 0]

i = 0: [3, 0, 0]

i = 1: [3, 28, 0]

i = 2: [3, 28, 224]

# Memory Aliasing

Must consider possibility that these updates will affect program behavior

```
/* Sum rows is of n X n matrix a  
and store in vector b */  
void sum_rows1(double *a, double *b, long n) {  
    long i, j;  
    for (i = 0; i < n; i++) {  
        b[i] = 0;  
        for (j = 0; j < n; j++)  
            b[i] += a[i*n + j];  
    }  
}
```

```
double A[9] =  
{ 0, 1, 2,  
 4, 8, 16,  
 32, 64, 128};  
double *B = A+3;  
sum_rows1(A, B, 3);
```

Value of B:

init: [4, 8, 16]

i = 0: [3, 8, 16]

i = 1: [3, 22, 16]

i = 2: [3, 22, 224]

3 0 16

3 22? What?

# Memory Aliasing

Must consider possibility that these updates will affect program behavior

```
/* Sum rows is of n X n matrix a  
and store in vector b */  
void sum_rows1(double *a, double *b, long n) {  
    long i, j;  
    for (i = 0; i < n; i++) {  
        b[i] = 0;  
        for (j = 0; j < n; j++)  
            b[i] += a[i*n + j];  
    }  
}
```

```
double A[9] =  
{ 0, 1, 2,  
 4, 8, 16,  
 32, 64, 128};  
double *B = A+3;  
sum_rows1(A, B, 3);
```

i=0: 4 8 16  
i=0: 3 8 16  
i=1: 3 0 16  
i=1: 3 3 16  
i=1: 3 6 16  
i=1: 3 22 16  
i=2: 3 22 224

# Removing Aliasing

No need to store intermediate results

```
/* Sum rows is of n X n matrix a  
and store in vector b */  
void sum_rows2(double *a, double *b, long n) {  
    long i, j;  
    for (i = 0; i < n; i++) {  
        double val = 0;  
        for (j = 0; j < n; j++)  
            val += a[i*n + j];  
        b[i] = val;  
    }  
}
```

```
double A[9] =  
{ 0, 1, 2,  
 4, 8, 16,  
 32, 64, 128};  
double *B = A+3;  
sum_rows1(A, B, 3);
```

Value of B:

init: [4, 8, 16]

i = 0: [3, 8, 16]

i = 1: [3, 27, 16]

i = 2: [3, 27, 224]

4 8 16  
3 8 16  
3 27 16  
3 27 224

# Optimization Blocker: Memory Aliasing

---

## Aliasing

- Two different memory references specify single location
- Easy to have happen in C
  - Since allowed to do address arithmetic
  - Direct access to storage structures
- Get in habit of introducing local variables
  - Accumulating within loops
  - Your way of telling compiler not to check for aliasing

# Summary

---

Must optimize at multiple levels:

- algorithm, data representations, procedures, and loops

Must understand system to optimize performance

Generally Useful Optimizations

- Code motion/precomputation
- Strength reduction
- Sharing of common subexpressions
- Removing unnecessary procedure calls

Optimization Blockers

- Procedure calls
- Memory aliasing