


OPTIMIZE MEMORY ACCESS

Jo, Heeseung

The CPU-Memory Gap

The increasing gap between DRAM, disk, and CPU speeds



Principle of Locality (1)

Temporal locality

- Recently referenced items are likely to be referenced in the near future

Spatial locality

- Items with nearby addresses tend to be referenced close together in time

Principle of Locality (2)

Locality example:

```
sum = 0;  
for (i = 0; i < n; i++)  
    sum += a[i];  
return sum;
```

Data

- Reference array elements in succession
- Reference sum each iteration

Spatial locality
Temporal locality

Instructions

- Reference instructions in sequence
- Cycle through loop repeatedly

Spatial locality
Temporal locality


Principle of Locality (3)

How to exploit temporal locality?


- Speed up data accesses by caching data in faster storage
- Caching in multiple levels: form a memory hierarchy:
 - The lower levels of the memory hierarchy tend to be slower, but larger and cheaper

How to exploit spatial locality?

- Larger cache line size
 - Cache nearby data together



Memory Hierarchy




Caching (1)

Cache

- A smaller, faster storage
- Improves the average access time
- Exploits both temporal and spatial locality

고속↓ 속도↑



Caching (2)

Cache performance metrics

- Average memory access time = $T_{\text{hit}} + R_{\text{miss}} * T_{\text{miss}}$
- Hit time (T_{hit})
 - Time to deliver a line in the cache to the processor
 - Includes time to determine whether the line is in the cache
 - 1 clock cycle for L1, 3 ~ 8 clock cycles for L2
- Miss rate (R_{miss})
 - Fraction of memory references not found in cache (misses/references)
 - 3 ~ 10% for L1, < 1% for L2
- Miss penalty (T_{miss})
 - Additional time required because of a miss
 - Typically 25 ~ 100 cycles for main memory

Caching (3)

Cache design issues

- Cache size
 - 8KB ~ 64KB for L1
- Cache line size
 - Typically, 32B or 64B for L1
- Lookup
 - Fully associative
 - Set associative: 2-way, 4-way, 8-way, 16-way, etc.
 - Direct mapped
- Replacement
 - LRU (Least Recently Used)
 - FIFO (First-In First-Out), etc.

Matrix Multiplication (1)

Description

- Multiply $N \times N$ matrices
- $O(N^3)$ total operations

$$(i, *) \begin{array}{|c|} \hline \text{---} \\ \hline \end{array} \times \begin{array}{|c|} \hline \text{---} \\ \hline \end{array} = \begin{array}{|c|} \hline \text{---} \\ \hline \end{array}$$

A B C

$(*, j)$ (i, j)

*Variable sum
held in register*

```
/* ijk */
for (i=0; i<n; i++) {
    for (j=0; j<n; j++) {
        sum = 0.0;
        for (k=0; k<n; k++)
            sum += a[i][k] * b[k][j];
        c[i][j] = sum;
    }
}
```

Assumptions


- Line size = 32B (big enough for 4 64-bit words)
- Matrix dimension (N) is very large
- Cache is not big enough to hold multiple rows

Matrix Multiplication (2)

Matrix multiplication (ijk)

```
/* ijk */
for (i=0; i<n; i++) {
    for (j=0; j<n; j++) {
        sum = 0.0;
        for (k=0; k<n; k++)
            sum += a[i][k] * b[k][j];
        c[i][j] = sum;
    }
}
```

Inner loop:



Misses per Inner Loop Iteration:

A
0.25

B
1.0


C
0.0

Matrix Multiplication (3)

Matrix multiplication (jik)

```
/* jik */  
for (j=0; j<n; j++) {  
    for (i=0; i<n; i++) {  
        sum = 0.0;  
        for (k=0; k<n; k++)  
            sum += a[i][k] * b[k][j];  
        c[i][j] = sum  
    }  
}
```

Inner loop:



Misses per Inner Loop Iteration:

A
0.25

B
1.0


C
0.0

Matrix Multiplication (4)


Matrix multiplication (kij)

```
/* kij */  
for (k=0; k<n; k++) {  
    for (i=0; i<n; i++) {  
        r = a[i][k];  
        for (j=0; j<n; j++)  
            c[i][j] += r * b[k][j];  
    }  
}
```


Inner loop:



A



B



C

Fixed

Row-wise

Row-wise

Misses per Inner Loop Iteration:

A

0.0

B

0.25

C


0.25

Matrix Multiplication (5)


Matrix multiplication (ikj)

```
/* ikj */  
for (i=0; i<n; i++) {  
    for (k=0; k<n; k++) {  
        r = a[i][k];  
        for (j=0; j<n; j++)  
            c[i][j] += r * b[k][j];  
    }  
}
```


Inner loop:



A



B



C

Fixed

Row-wise Row-wise

Misses per Inner Loop Iteration:

A

0.0

B

0.25

C


0.25

Matrix Multiplication (6)

Matrix multiplication (jki)

```
/* jki */
for (j=0; j<n; j++) {
    for (k=0; k<n; k++) {
        r = b[k][j];
        for (i=0; i<n; i++)
            c[i][j] += a[i][k] * r;
    }
}
```

Inner loop:



Misses per Inner Loop Iteration:

A
1.0

B
0.0


C
1.0

Matrix Multiplication (7)

Matrix multiplication (kji)

```
/* kji */
for (k=0; k<n; k++) {
    for (j=0; j<n; j++) {
        r = b[k][j];
        for (i=0; i<n; i++)
            c[i][j] += a[i][k] * r;
    }
}
```

Inner loop:



Column - Fixed Column-
wise wise

Misses per Inner Loop Iteration:

A
1.0

B
0.0

C
1.0

Matrix Multiplication (8)

Summary

ijk (& jik):

- 2 loads, 0 stores
- misses/iter = 1.25

```
for (i=0; i<n; i++) {  
    for (j=0; j<n; j++) {  
        sum = 0.0;  
        for (k=0; k<n; k++)  
            sum += a[i][k] * [k][j];  
        c[i][j] = sum;  
    }  
}
```

kij (& ikj):

- 2 loads, 1 store
- misses/iter = 0.5

```
for (k=0; k<n; k++) {  
    for (i=0; i<n; i++) {  
        r = a[i][k];  
        for (j=0; j<n; j++)  
            c[i][j] += r * [k][j];  
    }  
}
```

jki (& kji):

- 2 loads, 1 store
- misses/iter = 2.0

```
for (j=0; j<n; j++) {  
    for (k=0; k<n; k++) {  
        r = b[k][j];  
        for (i=0; i<n; i++)  
            c[i][j] += a[i][k] * r;  
    }  
}
```

A
0.25

B
1.0

C
0.0

A
0.0

B
0.25

C
0.25

A
1.0

B
0.0

C
1.0

↑
miss ↓


↑
miss ↑

이전 가정 때문?: 2
1211 // : 1

Matrix Multiplication (9)

Performance in Pentium

- Miss rates are helpful but not perfect predictors
 - Code scheduling matters, too



Blocked Matrix Multiplication (1)

Improving temporal locality by blocking

- "Block" means a sub-block within the matrix
- Example: $N = 8$, sub-block size = 4

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \times \begin{bmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{bmatrix} = \begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix}$$

- Key idea: Sub-blocks (i.e., A_{xy}) can be treated just like scalars


$$C_{11} = A_{11}B_{11} + A_{12}B_{21} \quad C_{12} = A_{11}B_{12} + A_{12}B_{22}$$

$$C_{21} = A_{21}B_{11} + A_{22}B_{21} \quad C_{22} = A_{21}B_{12} + A_{22}B_{22}$$

Blocked Matrix Multiplication (2)

Performance in Pentium

- Blocking improves performance by a factor of two over unblocked version (ijk and jik)



Observations

Programmer can optimize for cache performance

- How data structures are organized
- How data are accessed
 - Nested loop structure
 - Blocking is a general technique

All systems favor "cache friendly code"

- Getting absolute optimum performance is very platform specific
 - Cache sizes, line sizes, associativities, etc.
- Can get most of the advantage with generic code
 - Keep working set reasonably small (temporal locality)
 - Use small strides (spatial locality)

지식채널 적절한 기술

- <https://www.youtube.com/watch?v=op4jznLfQQE>