

Processes

Jo, Heeseung

Today's Topics

What is the process?

How to implement processes?

What Is The Process?

Program?

vs.

Process?

vs.

Processor? CPU

vs.

Task? Job?

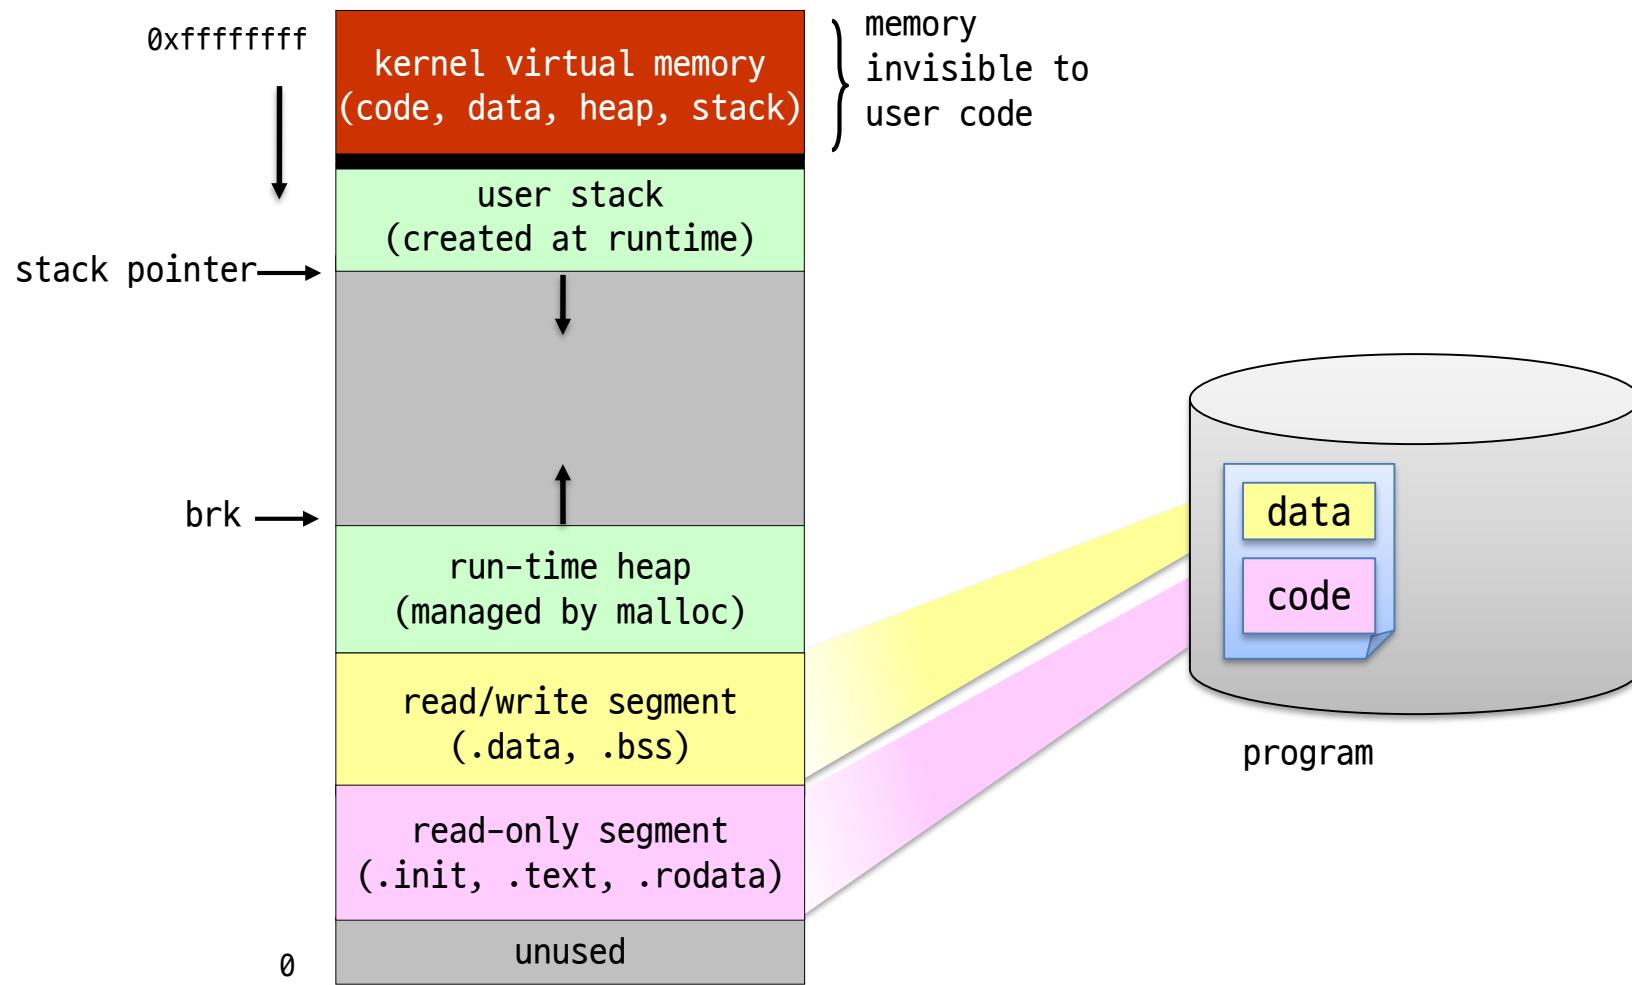
Process Concept (1)

What is the process?

- An instance of a program in execution
- An encapsulation of the flow of control in a program
- A dynamic and active entity
- The basic unit of execution and scheduling
- A process is named using its process ID (PID)
- A process includes:
 - CPU contexts (registers)
 - OS resources (memory, open files, etc.)
 - Other information (PID, state, owner, etc.)

Process Concept (2)

Process in memory

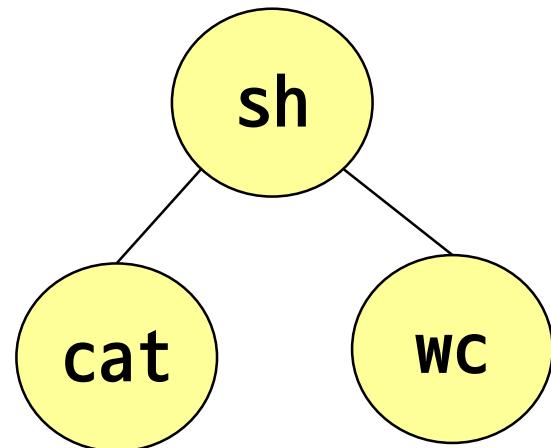


Process Creation (1)

Process hierarchy

- One process can create another process: parent-child relationship
- UNIX calls the hierarchy a "process group"
- Windows has no concept of process hierarchy
- Browsing a list of processes:
 - ps in UNIX
 - taskmgr (Task Manager) in Windows

```
$ cat file1 | wc
```



Process Creation (2)

Process creation events

- Calling a system call
 - `fork()` in POSIX, `CreateProcess()` in Win32
 - Shells or GUIs use this system call internally
- System initialization
 - *init* process
 - PID 1 process

Process Creation (3)

Resource sharing

- Parent may inherit **all or a part of resources** and **privileges** for its children
 - UNIX: User ID, open files, etc.

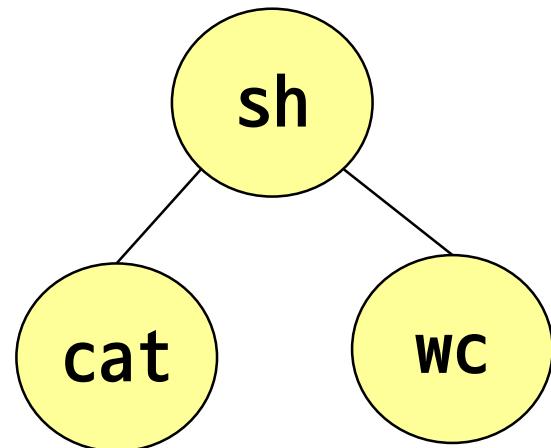
Execution

- Parent may either wait for it to finish, or it may continue in parallel

Address space

- Child duplicates the parent's address space or has a program loaded into it

```
$ cat file1 | wc
```



Process Termination

Process termination events

- Normal exit (voluntary)
- Error exit (voluntary)
- Fatal error (involuntary)
 - Exceed allocated resources
 - Segmentation fault
 - Protection fault, etc.
- Killed by another process (involuntary)
 - By receiving a signal

```
#include <stdio.h>

int main()
{
    int i, fd;
    char buf[100];

    fd=open("a.txt", "r");
    if (fd==NULL)
        return -1;
    read(fd, buf, 1000);

    return 0;
}
```

fork()

fork() system call

- Creating a child process
 - Copy the whole virtual address space of parent to create a child process
 - Copy internal data structures to manage a child process
-
- Parent get the pid of a child
 - Child get 0 value

fork()

```
#include <sys/types.h>
#include <unistd.h>

int main()
{
    int pid;

    pid = fork();
    if (pid == 0)
        /* child */
        printf ("Child of %d is %d\n",
               getppid(), getpid());
    else
        /* parent */
        printf ("I am %d. My child is %d\n",
               getpid(), pid);
}
```

```
#include <sys/types.h>
#include <unistd.h>

int main()
{
    int pid;

    pid = fork();
    if (pid == 0)
        /* child */
        printf ("Child of %d is %d\n",
               getppid(), getpid());
    else
        /* parent */
        printf ("I am %d. My child is %d\n",
               getpid(), pid);
}
```

fork(): Example Output

```
% ./a.out
```

I am 30000. My child is 30001.

Child of 30000 is 30001.

```
% ./a.out
```

Child of 30002 is 30003.

I am 30002. My child is 30003.

```
#include <sys/types.h>
#include <unistd.h>

int main()
{
    int pid;

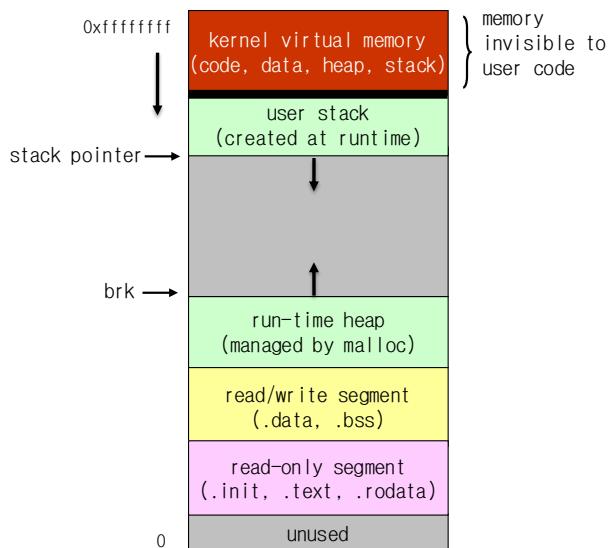
    pid = fork();
    if (pid == 0)
        /* child */
        printf ("Child of %d is %d\n",
               getppid(), getpid());
    else
        /* parent */
        printf ("I am %d. My child is %d\n",
               getpid(), pid);
}
```

fork() and Virtual Address Space

```
#include <sys/types.h>
#include <unistd.h>

int main()
{
    int pid;

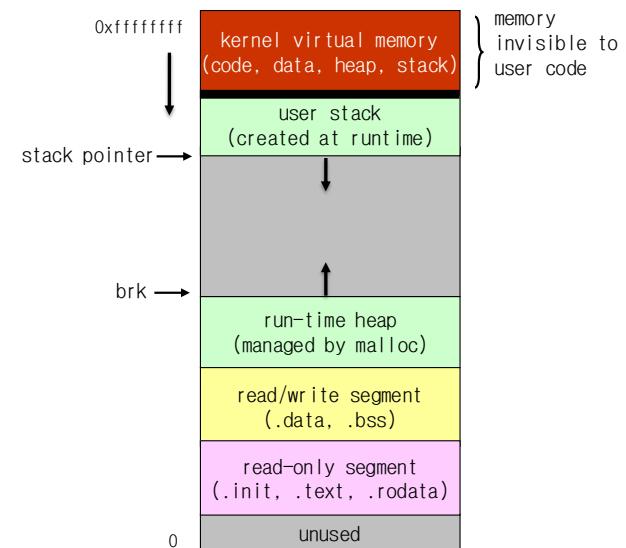
    pid = fork();
    if (pid == 0)
        /* child */
        printf ("Child of %d is %d\n",
               getppid(), getpid());
    else
        /* parent */
        printf ("I am %d. My child is %d\n",
               getpid(), pid);
}
```



```
#include <sys/types.h>
#include <unistd.h>

int main()
{
    int pid;

    pid = fork();
    if (pid == 0)
        /* child */
        printf ("Child of %d is %d\n",
               getppid(), getpid());
    else
        /* parent */
        printf ("I am %d. My child is %d\n",
               getpid(), pid);
}
```

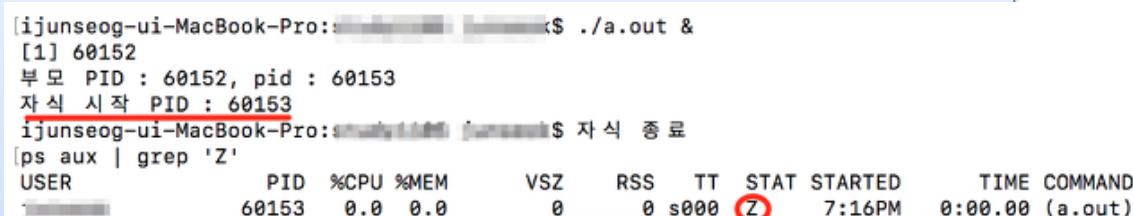


Zombie vs. orphan process

Zombie process (defunct process)

- A process that completed execution (via the exit system call) but still has an entry in the process table
- This occurs for the child processes, where the entry is still needed to allow the parent process to read its child's exit status

```
int main() {  
    pid_t childPid;  
  
    childPid = fork();  
  
    if (childPid > 0) { // parent process  
        printf("parent PID : %ld, pid : %d\n", (long)getpid(), childPid);  
        sleep(30);  
        printf("parent exit\n");  
        exit(0);  
    }  
    else if (childPid == 0){ // 자식 코드  
        printf("child PID : %ld\n", (long)getpid());  
        sleep(1);  
        printf("child exit\n");  
        exit(0);  
    }  
    return 0;  
}
```



The terminal window shows the execution of a C program named 'a.out'. The command is ./a.out &. The output includes the parent's PID (60152), the child's PID (60153), and the command ps aux | grep 'Z' which lists the child process as a zombie (Z state).

USER	PID	%CPU	%MEM	VSZ	RSS	TT	STAT	STARTED	TIME	COMMAND
[redacted]	60153	0.0	0.0	0	0	s000	Z	7:16PM	0:00.00	(a.out)

Zombie vs. orphan process

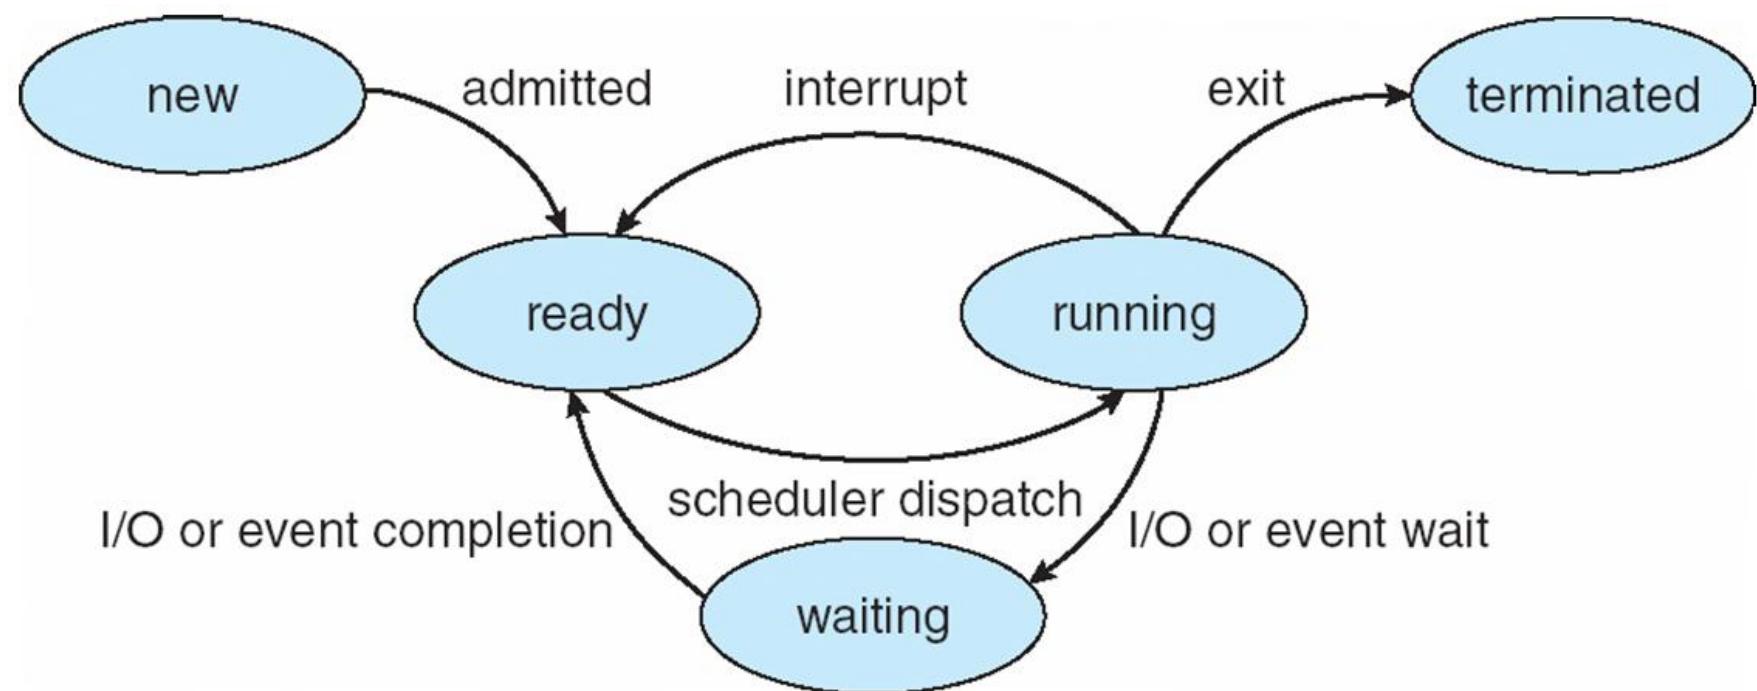
Orphan process

- A process whose parent process has finished or terminated, though it remains running itself
- Any orphaned process will be immediately adopted by the special init system process

```
int main() {  
  
    pid_t childPid;  
    int i;  
  
    childPid = fork();  
  
    if (childPid > 0) { // parent process  
        printf("parent PID : %ld, pid : %d\n", (long)getpid(), childPid);  
        sleep(2);  
        printf("parent exit\n");  
        exit(0);  
    }  
    else if (childPid == 0){ // child process  
        for(i=0;i<10;i++) {  
            printf("child PID : %ld parent PID : %ld\n", (long)getpid(), (long)getppid());  
            sleep(1);  
        }  
        printf("child exit\n");  
        exit(0);  
    }  
}
```

```
[ijunseog-ui-MacBook-Pro: ~] $ ./a.out  
부모 PID : 46797, pid : 46798  
자식 시작  
자식 PID : 46798 부모 PID : 46797  
자식 PID : 46798 부모 PID : 46797  
자식 PID : 46798 부모 PID : 46797  
부모 종료  
[ijunseog-ui-MacBook-Pro: ~] $ 자식 PID : 46798 부모 PID : 1  
자식 종료
```

Process State Transition (1)



Process State Transition (2)

Linux example

```
root      991  1  0 Mar19 ?      Ss  0:03 nmbd -D
root     1019  934  0 Mar19 ?      S   0:00 smbd -F
root     1038  2  0 Mar19 ?      S   0:00 [hd-audio1]
root     1039  1  0 Mar19 tty4    Ss+ 0:00 /sbin/getty -8 38400 tty4
root     1043  1  0 Mar19 tty5    Ss+ 0:00 /sbin/getty -8 38400 tty5
root     1051  1  0 Mar19 tty2    Ss+ 0:00 /sbin/getty -8 38400 tty2
root     1052  1  0 Mar19 tty3    Ss+ 0:00 /sbin/getty -8 38400 tty3
root     1057  1  0 Mar19 tty6    Ss+ 0:00 /sbin/getty -8 38400 tty6
root     1063  1  0 Mar19 ?      Ss   0:00 /usr/sbin/irqbalance
root    1130  1  0 Mar19 ?      Ss   0:00 cron
                                         Ss   0:00 atd
mysql   1186  1  0 Mar19 ?      Ssl  0:13 /usr/sbin/mysqld
root    1232  1  0 Mar19 ?      SI   0:00 /usr/sbin/console-kit-daemon --no-daemon
root    1296  1  0 Mar19 ?      Ss   0:01 /usr/sbin/apache2 -k start
root    1380  1  0 Mar19 tty1   Ss+  0:00 /sbin/getty -8 38400 tty1
root    2840  401 0 Mar19 ?     S<   0:00 udevd --daemon
root    2930  401 0 Mar19 ?     S<   0:00 udevd --daemon
root    3011  2  0 Mar19 ?      S   0:00 [kdmflush]
www-data 3133 1296 0 Mar19 ?    S   0:00 /usr/sbin/apache2 -k start
root    5450  2  0 10:09 ?      S   0:00 [flush-8:0]
www-data 6470 1296 0 12:03 ?    S   0:00 /usr/sbin/apache2 -k start
www-data 6471 1296 0 12:03 ?    S   0:00 /usr/sbin/apache2 -k start
www-data 6580 1296 0 13:20 ?    S   0:00 /usr/sbin/apache2 -k start
www-data 6581 1296 0 13:20 ?    S   0:00 /usr/sbin/apache2 -k start
www-data 6584 1296 0 13:22 ?    S   0:00 /usr/sbin/apache2 -k start
www-data 7133 1296 0 14:16 ?    S   0:00 /usr/sbin/apache2 -k start
root    7142  953 0 14:24 ?    Ss   0:00 sshd: root@pts/0
root    7205  7142 0 14:24 pts/0  Ss   0:00 -bash
root    7261  7205 0 14:27 pts/0  R+   0:00 ps -ef ax
```

R: Runnable

S: Sleeping

T: Traced or
Stopped

D: Uninterruptible
Sleep

Z: Zombie

<: High-priority task

N: Low-priority task

S: Session leader

+: In the foreground
process group

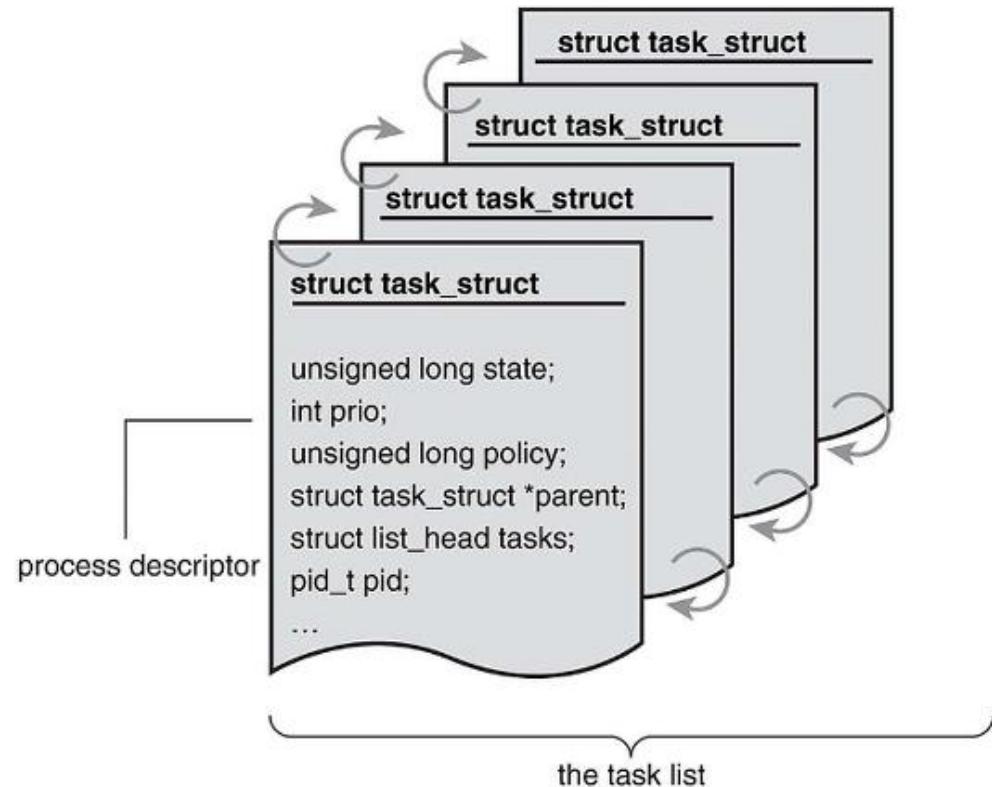
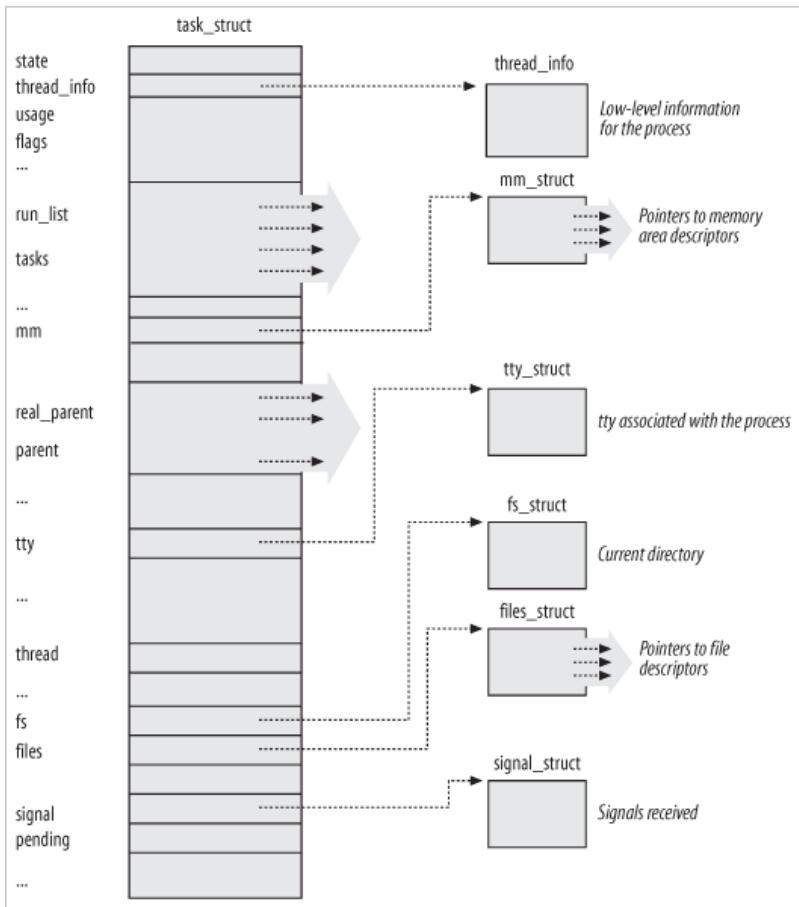
l: Multi-threaded

Process Data Structures

PCB (Process Control Block)

- Each PCB represents a process
- Contains all of the information about a process
 - Process state
 - Program counter
 - CPU registers
 - CPU scheduling information
 - Memory management information
 - Accounting information
 - I/O status information, etc.
- task_struct in Linux
 - 1456 bytes as of Linux 2.4.18
 - `include/linux/sched.h`

task_struct



Process Control Block (PCB)

Process management	Memory management	File management
Registers Program counter Program status word Stack pointer Process state Priority Scheduling parameters Process ID Parent process Process group Signals Time when process started CPU time used Children's CPU time Time of next alarm	Pointer to text segment Pointer to data segment Pointer to stack segment	Root directory Working directory File descriptors User ID Group ID

PCBs and Hardware State

When a process is running:

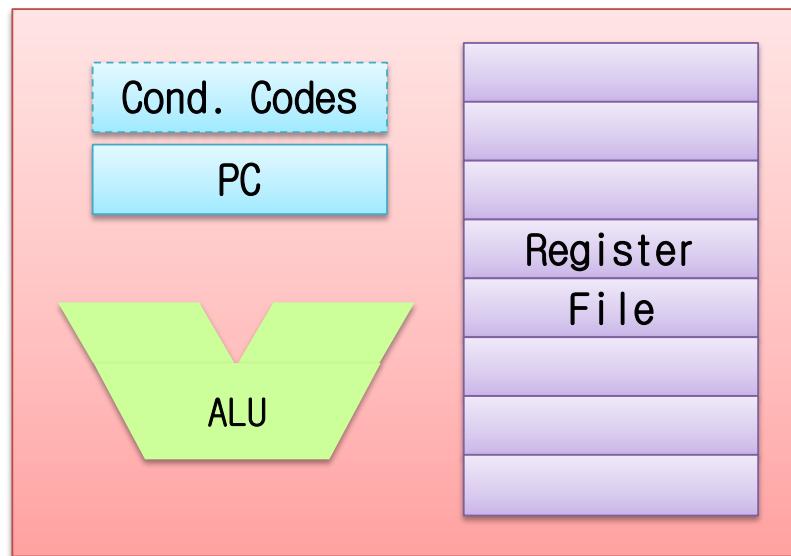
- Its hardware state is inside the CPU: PC, SP, registers

When the OS stops running a process:

- It saves the registers' values in the PCB

When the OS puts the process in the running state:

- It loads the hardware registers from the values in that process' PCB



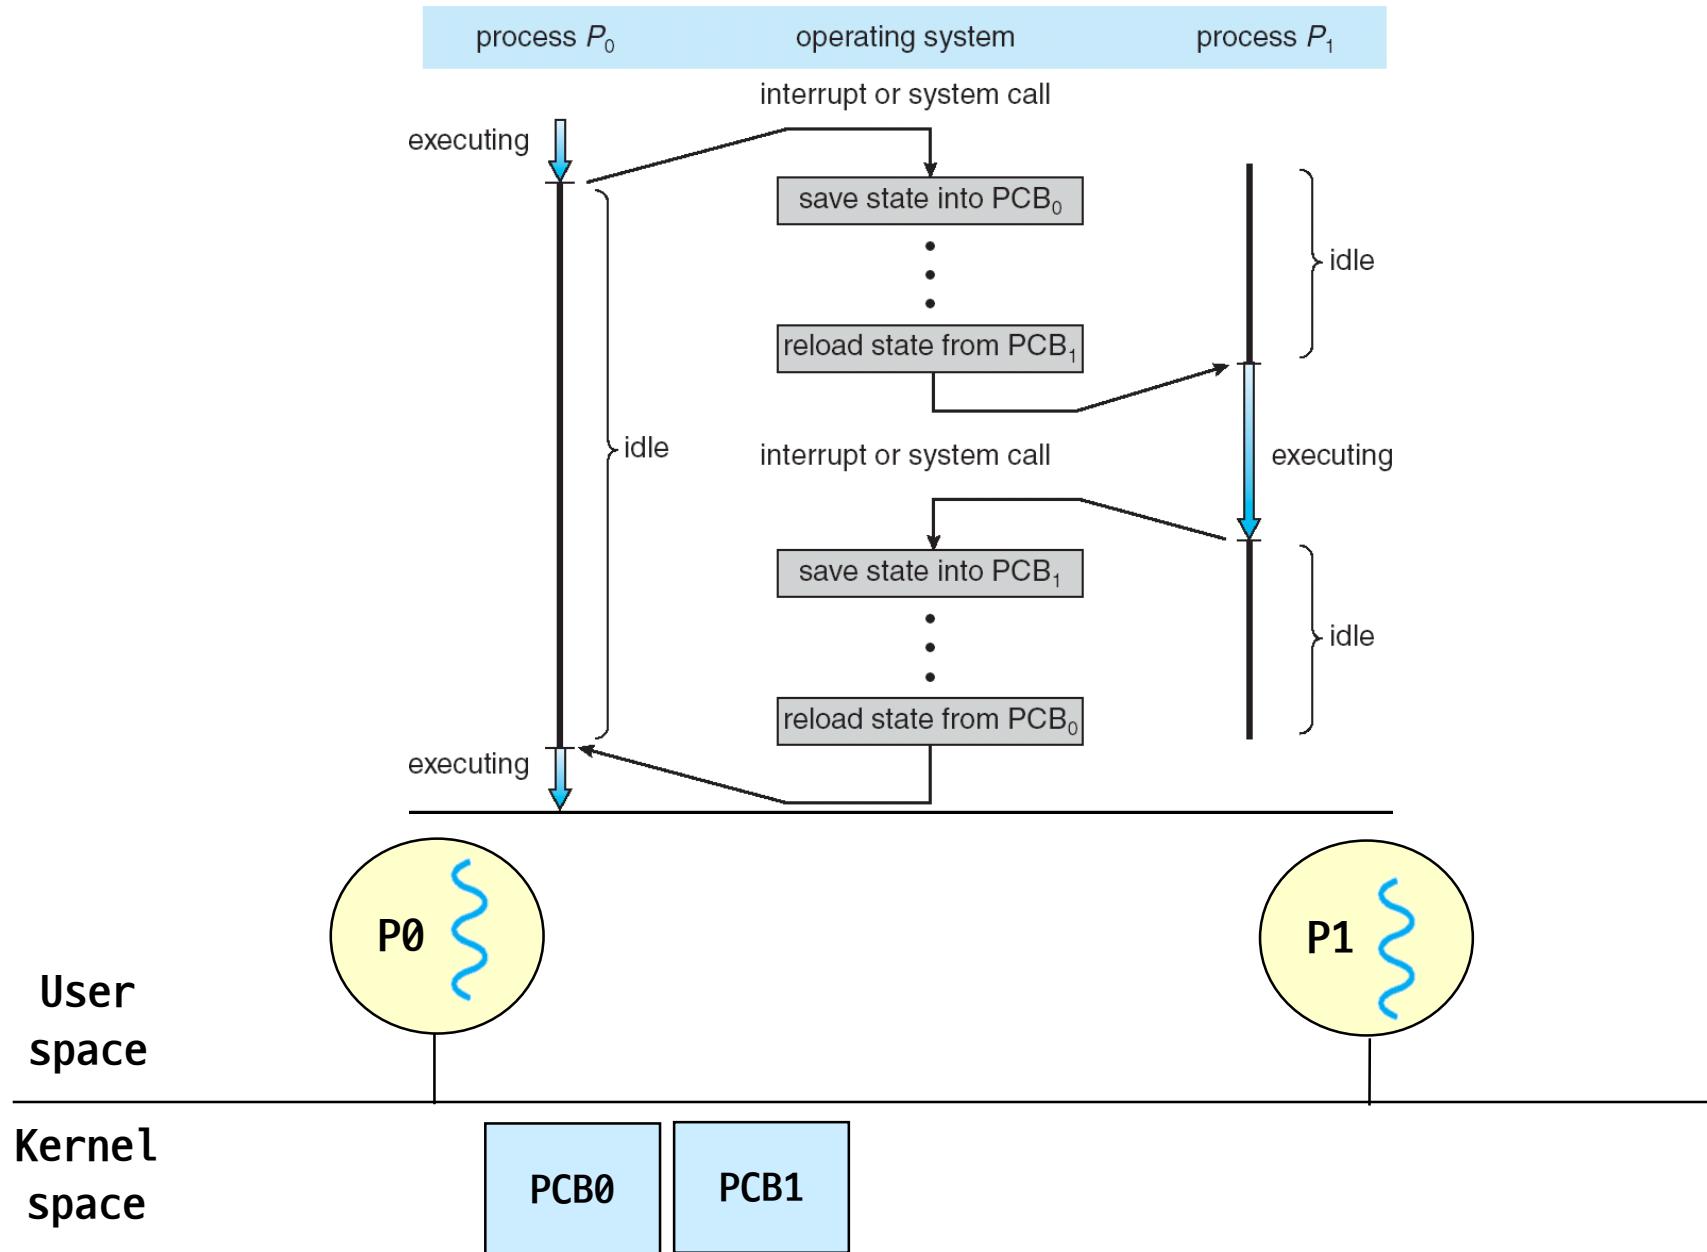
Process management	Memory management	File management
Registers Program counter Program status word Stack pointer Process state Priority Scheduling parameters Process ID Parent process Process group Signals Time when process started CPU time used Children's CPU time Time of next alarm	Pointer to text segment Pointer to data segment Pointer to stack segment	Root directory Working directory File descriptors User ID Group ID

Context Switch (1)

Context switch (or process switch)

- The act of **switching** the CPU from one process to another
- **Administrative overhead**
 - Saving and loading registers and memory maps
 - Flushing and reloading the memory cache
 - Updating various tables and lists, etc.
- Context switch overhead is dependent on hardware support
 - Multiple register sets in UltraSPARC
 - Advanced memory management techniques may require extra data to be switched with each context
- 100s or 1000s of switches/s typically

Context Switch (2)



Context Switch (3)

Linux example

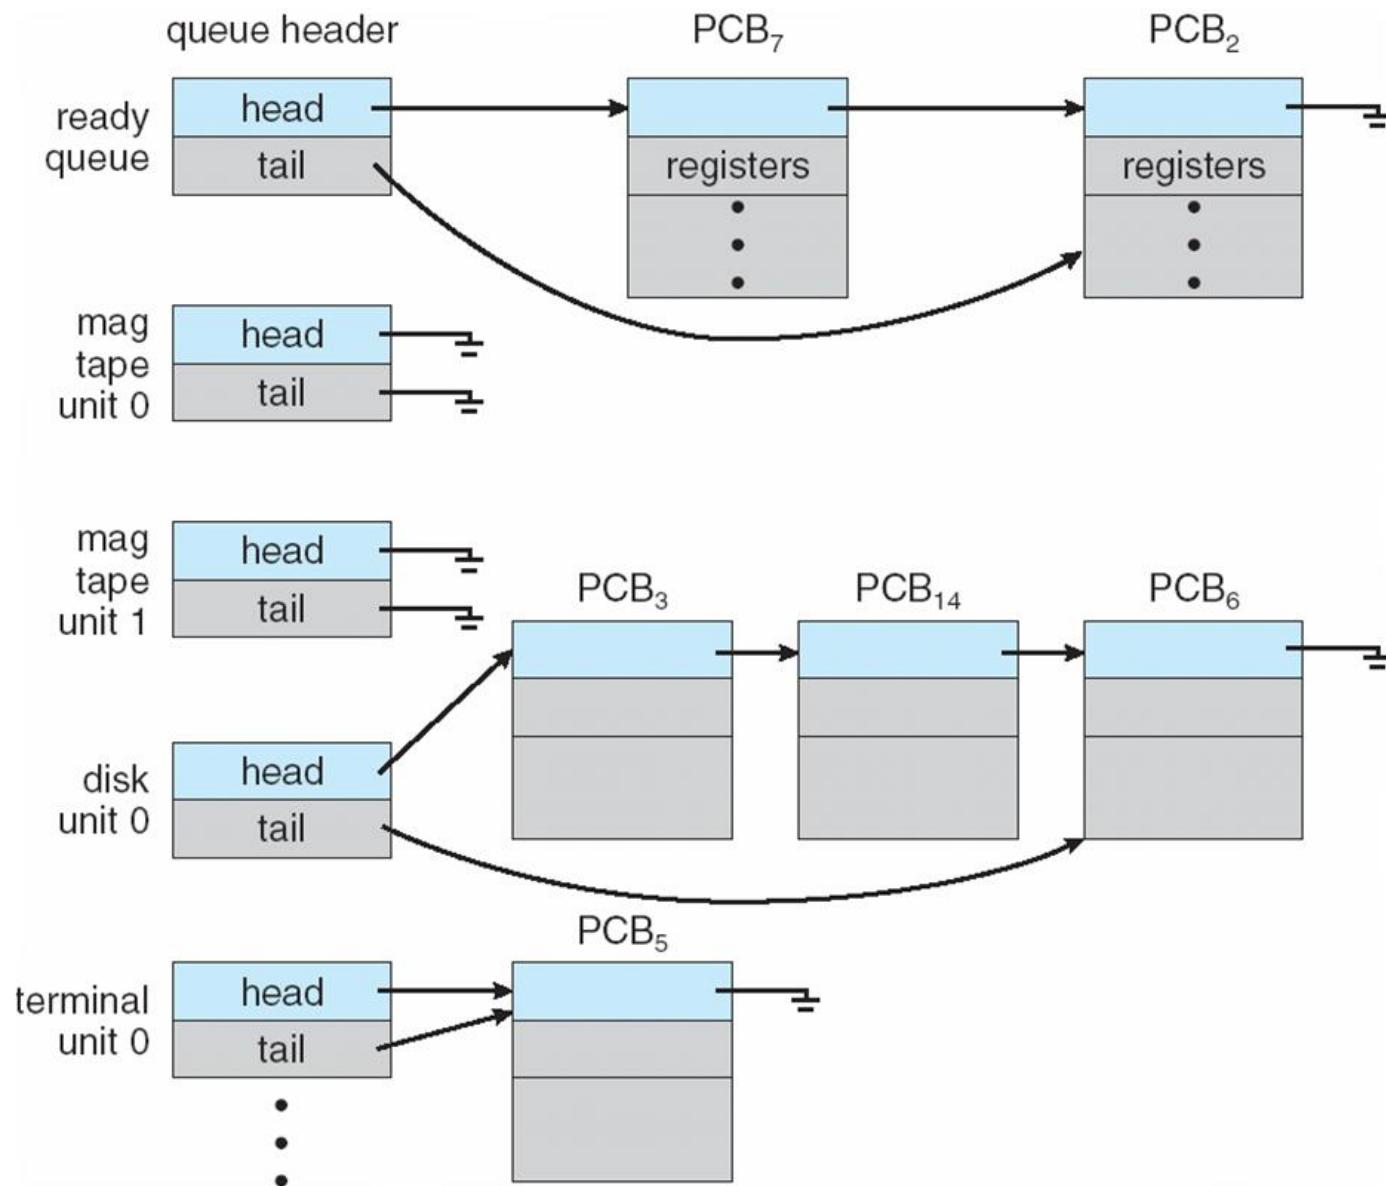
- Total 541,514,896 user ticks = 1511 hours = 63.0 days
 - Total 930,566,190 context switches
 - Roughly 86 context switches / sec (per CPU)

Process State Queues (1)

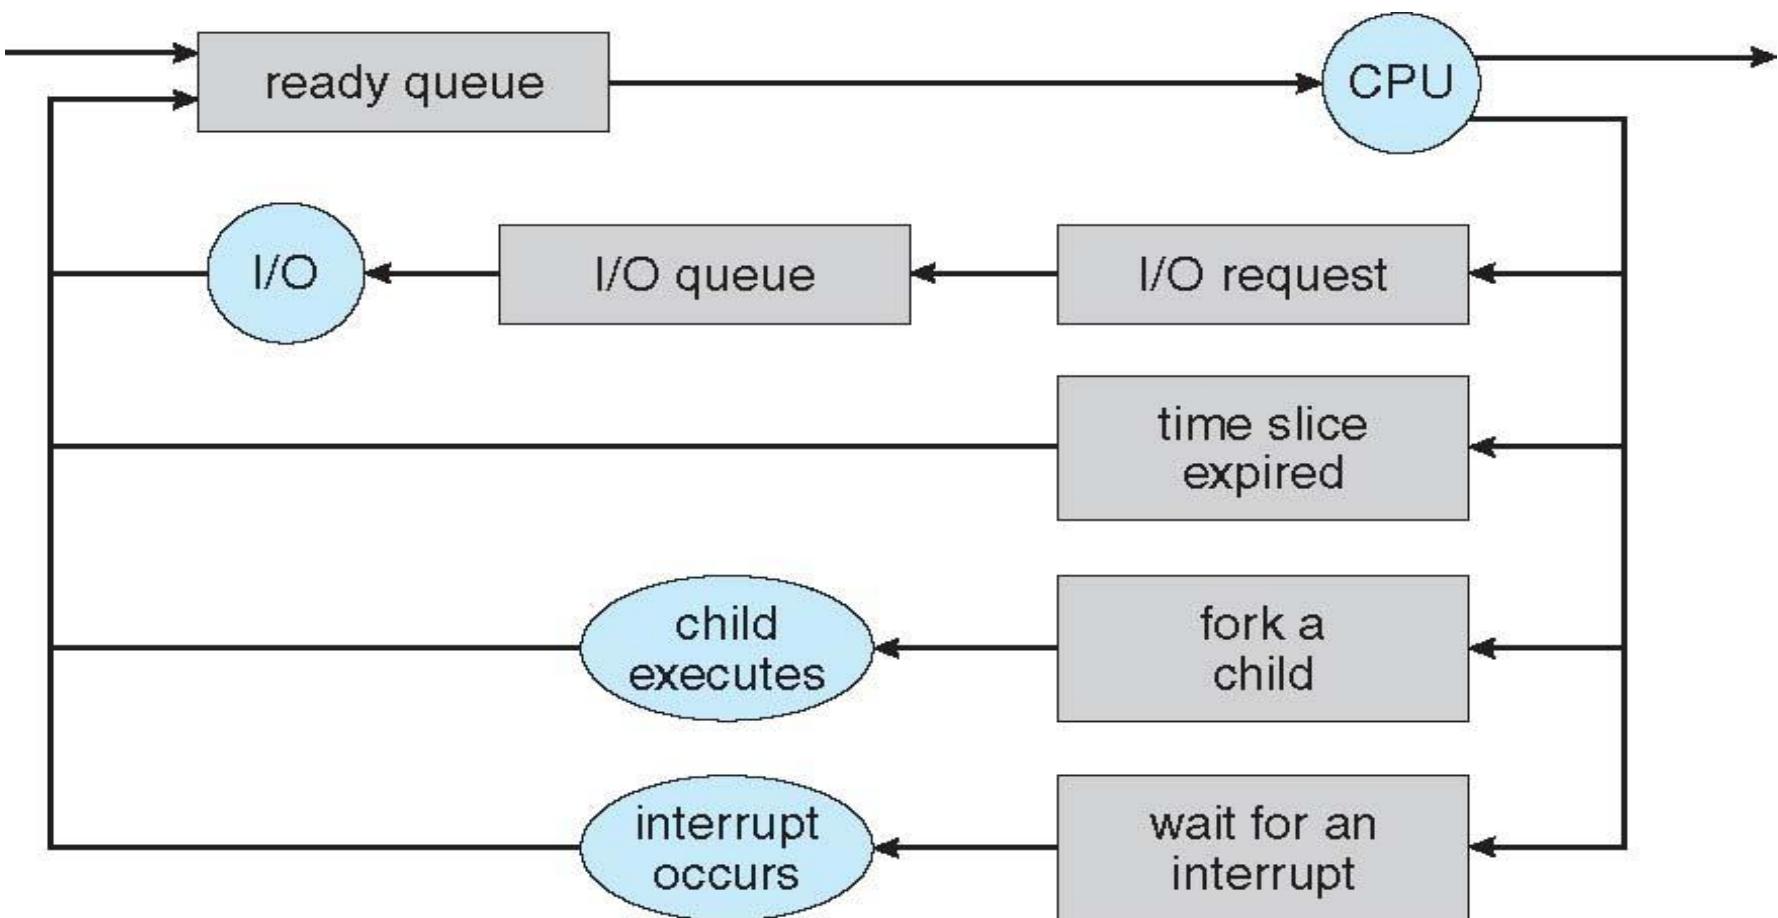
State queues

- The OS maintains a collection of queues that represent the state of all processes in the system
 - Ready queue
 - Wait queue(s): there may be many wait queues, one for each type of wait (device, timer, message, ...)
- Each PCB is queued onto a state queue according to its current state
- As a process changes state, its PCB is migrated between the various queues

Process State Queues (2)



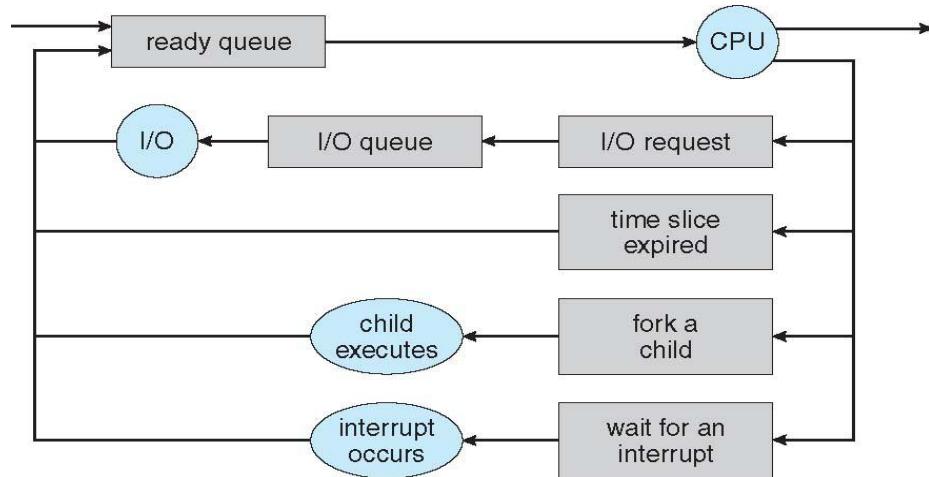
Process State Queues (3)



Process State Queues (4)

PCBs and state queues

- PCBs are data structures
 - Dynamically allocated inside OS memory
- When a process is created:
 - OS allocates a PCB for it
 - OS initializes PCB
 - OS puts PCB on the correct queue
- As a process computes:
 - OS moves its PCB from queue to queue
- When a process is terminated:
 - OS deallocates its PCB



Process Creation: UNIX (1)

int fork()

fork()

- Creates and initializes a new PCB
- Creates and initializes a new address space
- Initializes the address space with a copy of the entire contents of the address space of the parent
- Initializes the kernel resources to point to the resources used by parent (e.g., open files)
- Places the PCB on the ready queue
- Returns the child's PID to the parent, and zero to the child

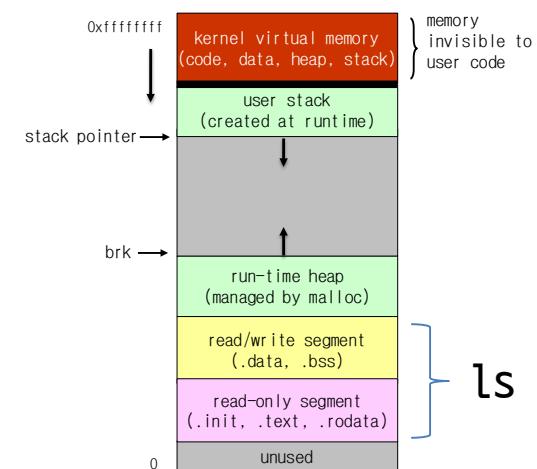
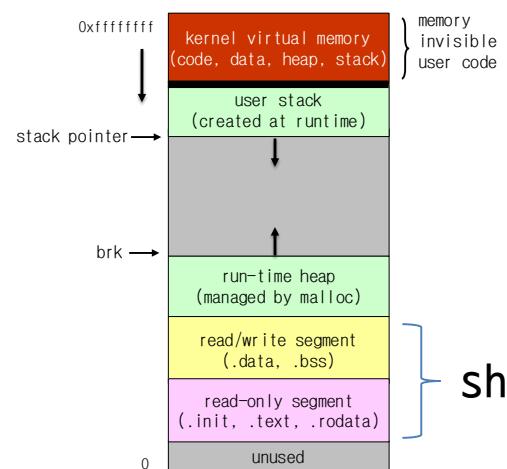
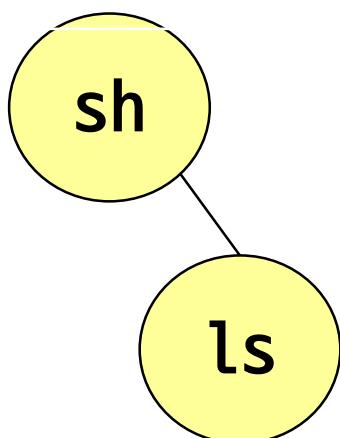
Process Creation: UNIX (2)

```
int exec (char *prog, char *argv[])
```

exec()

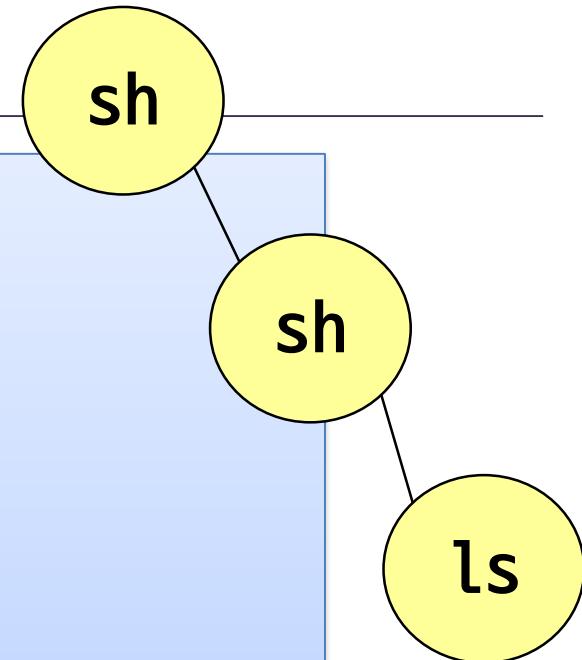
- Stops the current process
- Loads the program "prog" into the process' address space
- Initializes hardware context and args for the new program
- Places the PCB on the ready queue
 - Note: exec() does not create a new process

```
$ ls -aFl
```



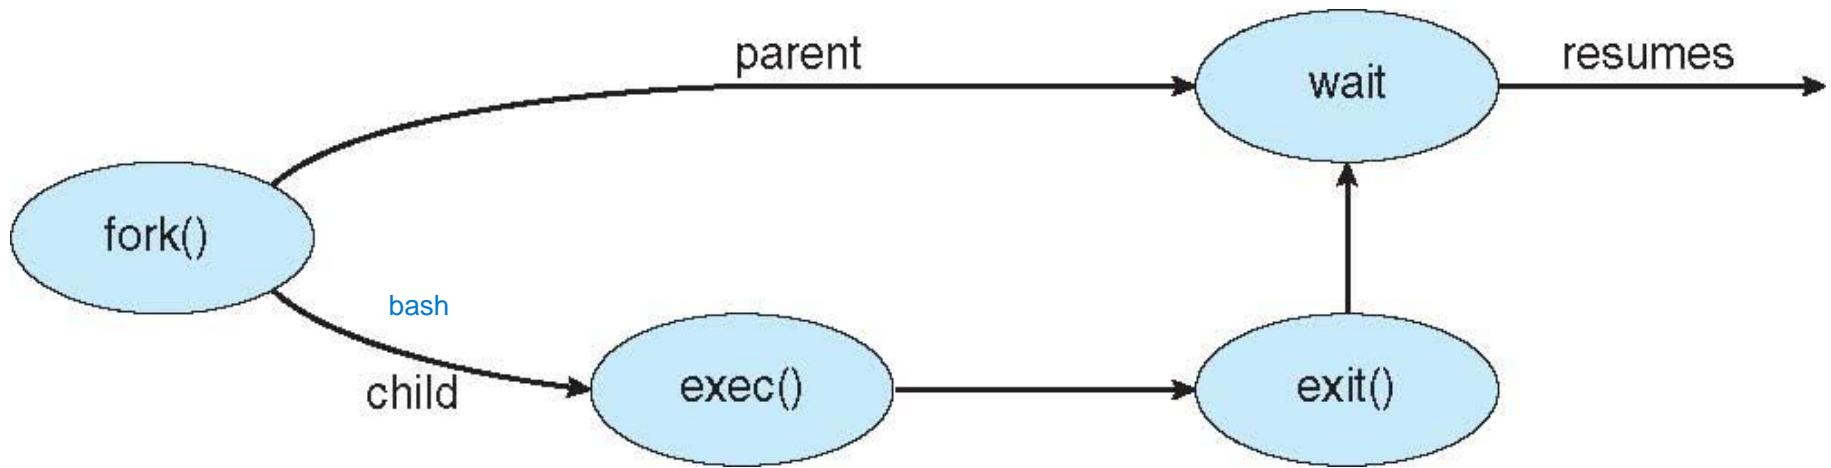
Simplified UNIX Shell

```
int main()
{
    while (1) {
        char *cmd = read_command();
        int pid;
        if ((pid = fork()) == 0) {
            /* Manipulate stdin/stdout/stderr for
               pipes and redirections, etc. */
            exec(cmd);
            panic("exec failed!");
        } else {
            wait (pid);
        }
    }
}
```



```
$ ls -aF
./  ../  gitclone.sh*  temp3/  temp3.tgz  token/
$
```

Simplified UNIX Shell



```
int main()
{
    while (1) {
        char *cmd = read_command();
        int pid;
        if ((pid = fork()) == 0) {
            /* Manipulate stdin/stdout/stderr for
               pipes and redirections, etc. */
            exec(cmd);
            panic("exec failed!");
        } else {
            wait (pid);
        }
    }
}
```

```
$ ls -aF
./ .. gitclone.sh* temp3/ temp3.tgz token/
$
```

Process Creation: NT

```
BOOL CreateProcess (char *prog, char *args, ...)
```

CreateProcess()

- Creates and initializes a new PCB
- Creates and initializes a new address space
- Loads the program specified by "prog" into the address space
- Copies "args" into memory allocated in address space
- Initializes the hardware context to start execution at main
- Places the PCB on the ready queue

Why fork()?

Very useful when the child ...

- Is cooperating with the parent
- Relies upon the parent's data to accomplish its task
- Example: Web server

```
While (1) {
    int sock = accept();
    if ((pid = fork()) == 0) {
        /* Handle client request */
    } else {
        /* Close socket */
    }
}
```

Inter-Process Communications

Inside a machine

- pipe
- FIFO
- Shared memory
- Sockets

Across machines

- Sockets
- RPCs (Remote Procedure Calls)
- Java RMI (Remote Method Invocation)