

# Threads

Jo, Heeseung

# Today's Topics

---

Why threads?

Threading issues

# Processes

---

## Heavy-weight

- A process includes many things:
  - An address space (all the code and data pages)
  - OS resources (e.g., open files) and accounting information
  - Hardware execution state (PC, SP, registers, etc.)
- Creating a new process is costly because all of the data structures must be allocated and initialized
  - Linux: over 100 fields in task\_struct (excluding page tables, etc.)
- Inter-process communication is costly, since it must usually go through the OS
  - Overhead of system calls and copying data

# Concurrent Servers: Processes

## Web server example

- Using `fork()` to create new processes to handle requests in parallel is overkill for such a simple task

```
while (1) {  
    int sock = accept();  
    if ((pid = fork()) == 0) {  
        /* Handle client request */  
    } else {  
        /* Close socket */  
    }  
}
```

# Cooperating Processes

---

## Example

- A web server, which forks off copies of itself to handle multiple simultaneous tasks
- Any parallel program on a multiprocessor

We need to:

- Create several processes that execute in parallel
- Cause each to map the same address space to share data
  - e.g., shared memory
- Have the OS schedule these processes in parallel

This is very inefficient!

- Space: PCB, page tables, etc.
- Time: creating OS structures, fork and copy address space, etc.

# Rethinking Processes

---

What's similar in these cooperating processes?

- They all use (share?) the same code and data (address space)
- They all use the same privilege
- They all use the same resources (files, sockets, etc.)

What's different?

- Each has its own hardware execution state:  
PC, registers, SP, and stack

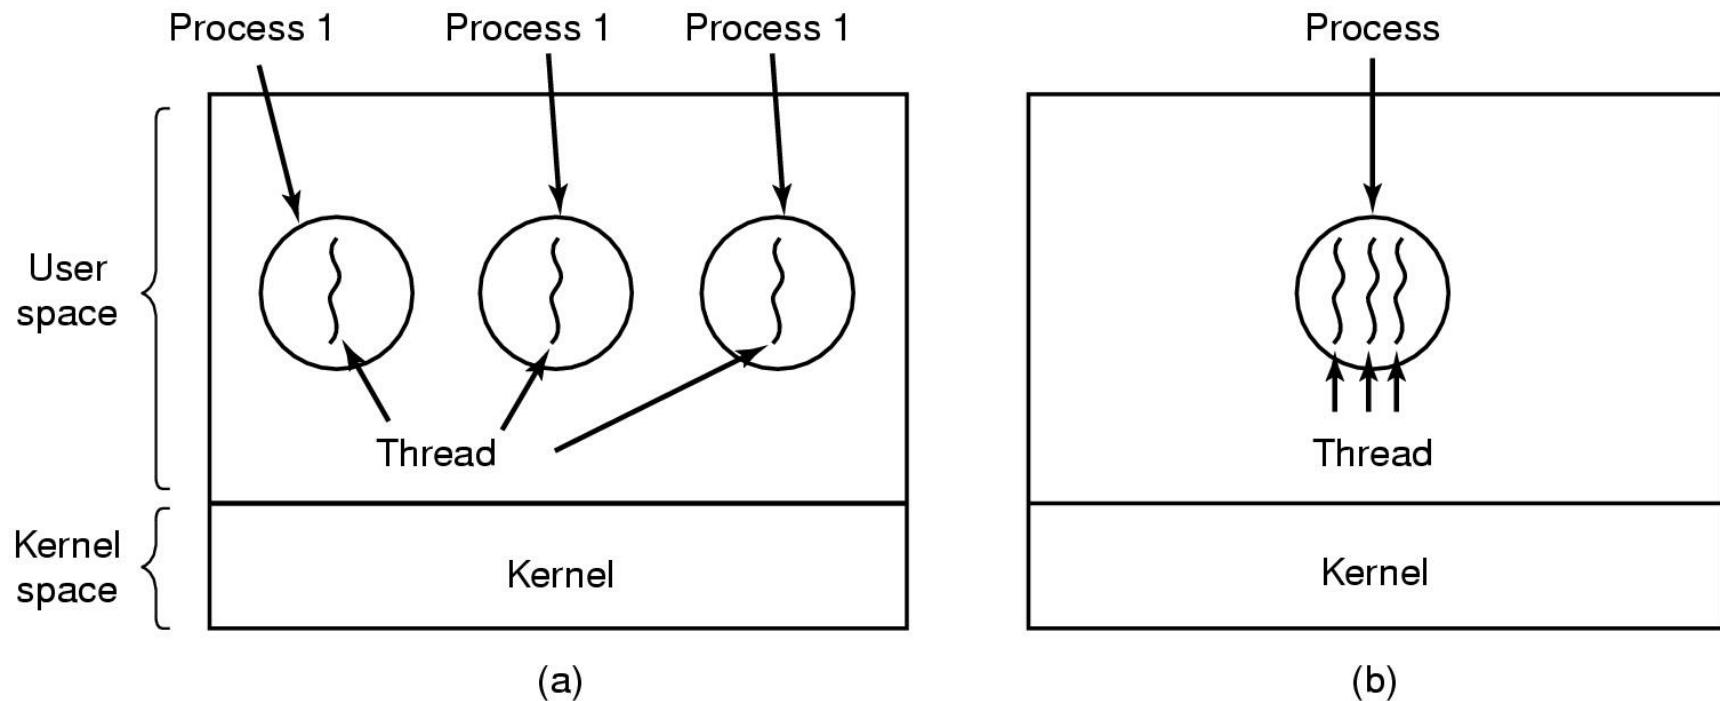
# Key Idea (1)

---

Separate the concept of a process from its execution state

- **Process**: address space, resources, other general process attributes
  - e.g., privileges
- **Execution state**: PC, SP, registers, etc.
- This execution state is usually called
  - Thread
  - Lightweight process (LWP)
  - Thread of control

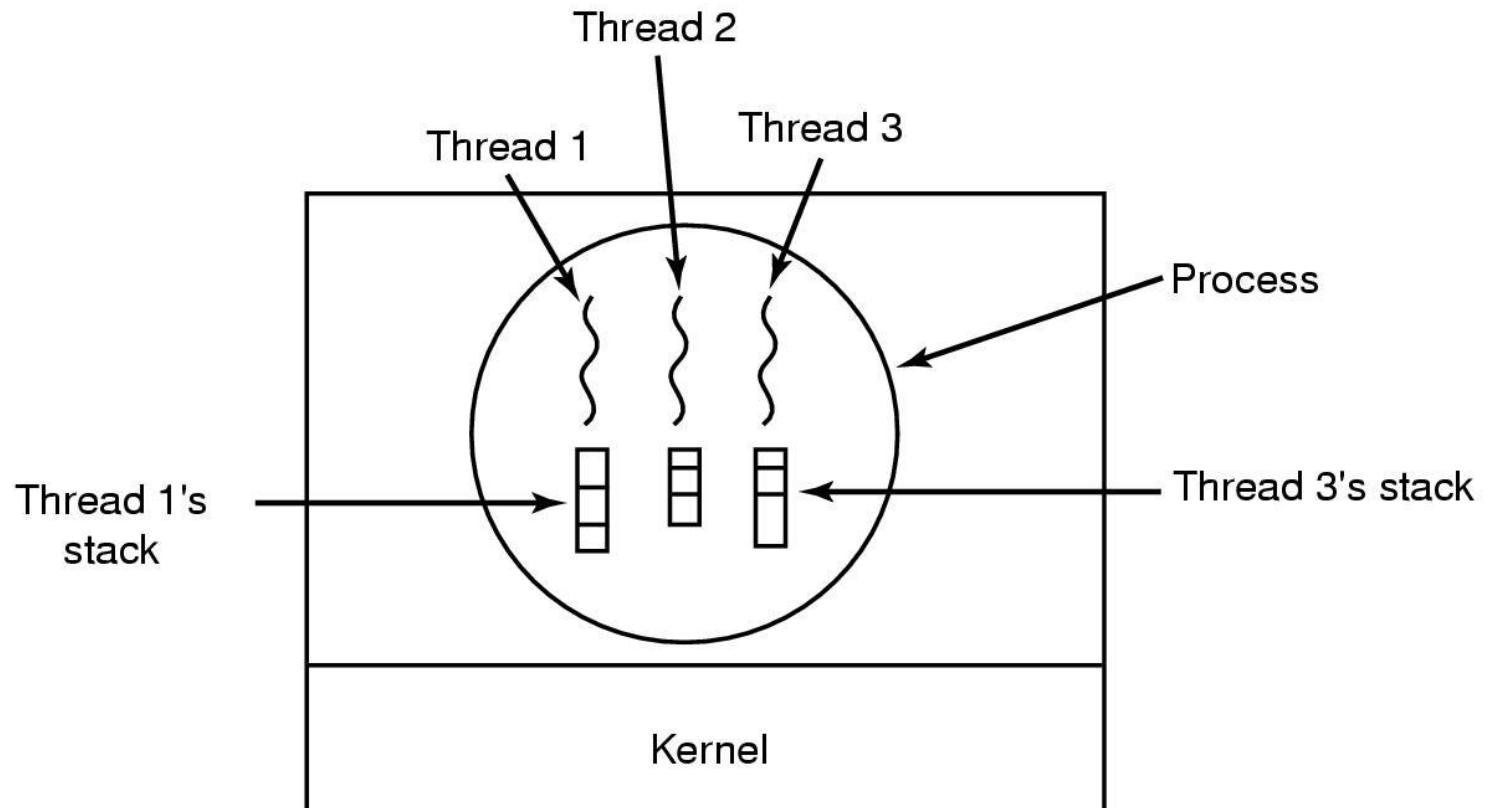
# Key Idea (2)



<b>Per process items</b>	<b>Per thread items</b>
Address space	Program counter
Global variables	Registers
Open files	Stack
Child processes	State
Pending alarms	
Signals and signal handlers	
Accounting information	

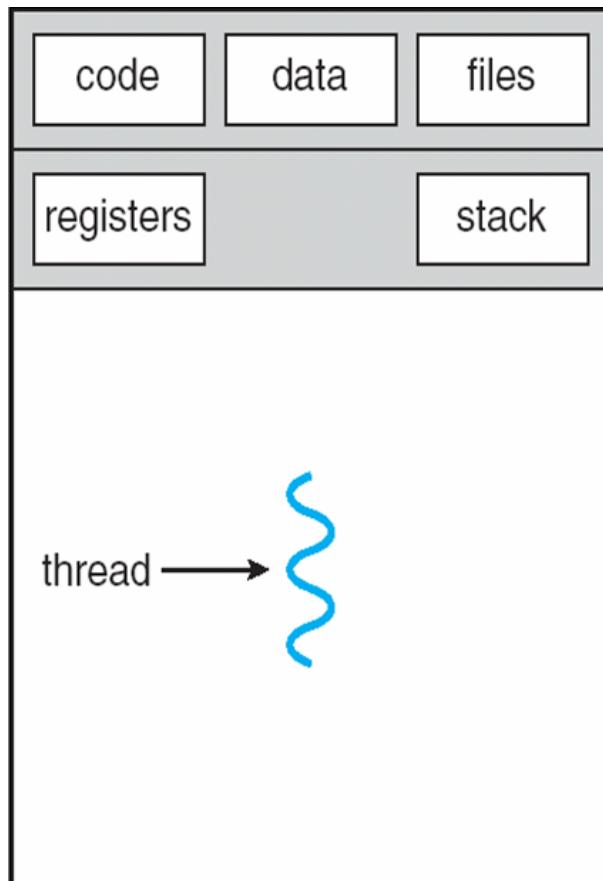
# Key Idea (3)

Each thread has its own stack

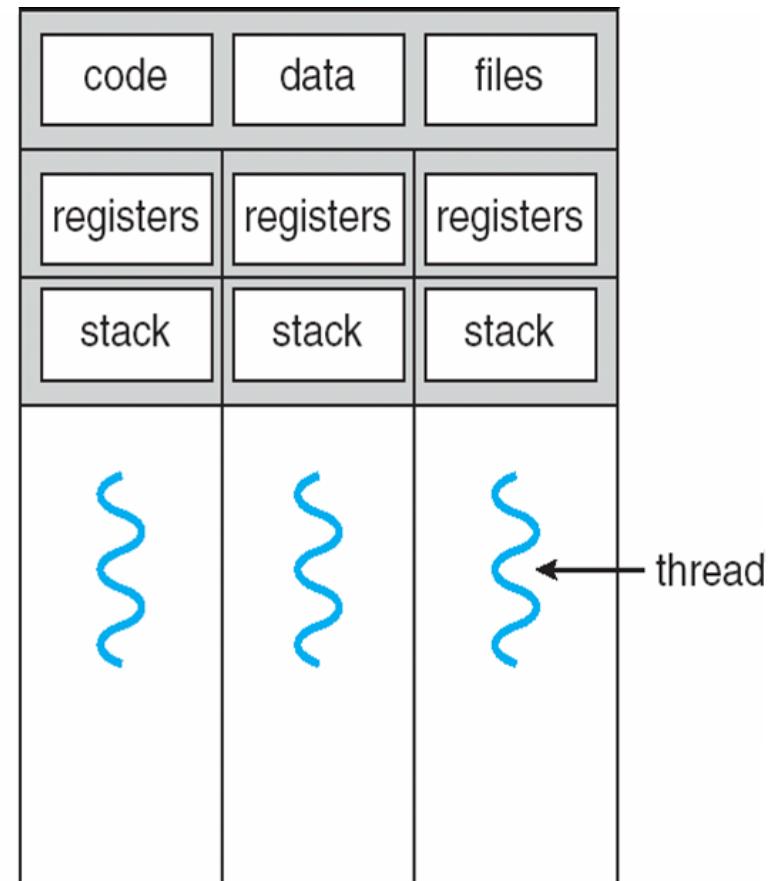


# Key Idea (4)

Each thread has its own stack



single-threaded process



multithreaded process

# What is a Thread?

---

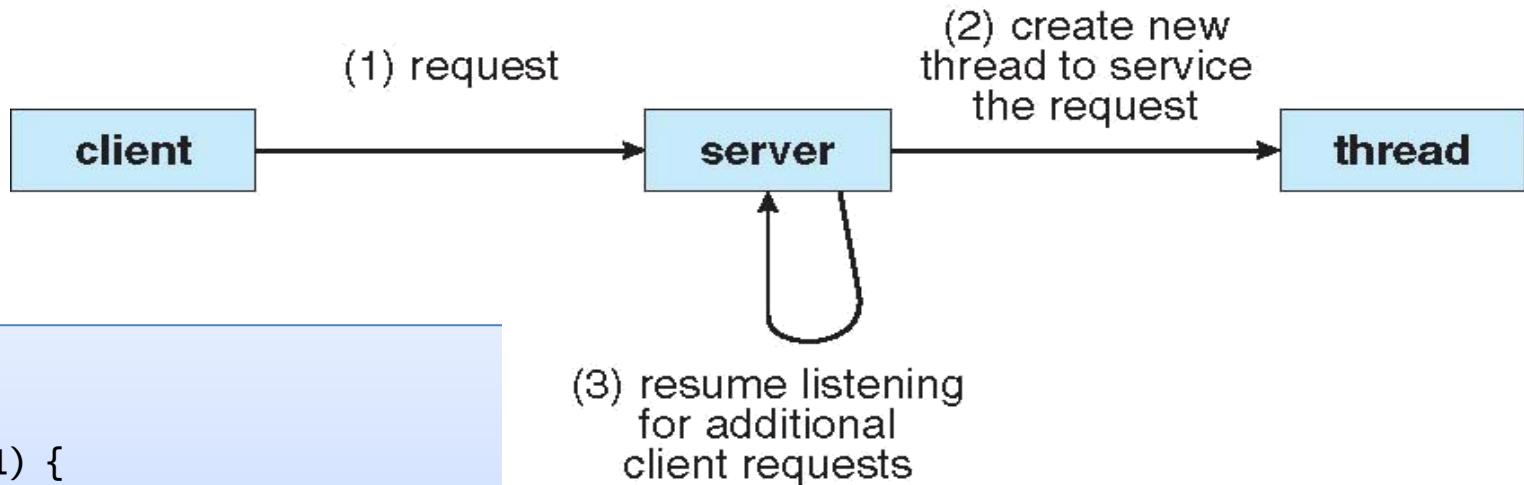
A thread of control (or a thread)

- A sequence of instructions being executed in a program
- Usually consists of
  - A program counter (PC), general registers
  - A stack to keep track of local variables and return addresses
- Threads share the process instructions and most of its data
  - A change in shared data by one thread can be seen by the other threads in the process
- Threads also share most of the OS state of a process

# Concurrent Servers: Threads

## Using threads

- We can create a new thread for each request



```
webserver ()  
{  
    while (1) {  
        int sock = accept();  
        create_thread (handle_request, sock);  
    }  
}  
handle_request (int sock)  
{  
    /* Process request */  
    close (sock);  
}
```

# Multithreading

---

## Benefits

- Creating concurrency is cheap
  - Time and memory consumption
- Improves program structure
- Higher throughput
  - By overlapping computation with I/O operations
- Better responsiveness (User interface / Server)
  - Can handle concurrent events (e.g., web servers)
- Better resource sharing
- Utilization of multiprocessor architectures
  - Allows building parallel programs

# Processes vs. Threads (1)

---

## Processes vs. Threads

- A thread is bound to a single process
- A process, however, can have multiple threads
- Sharing data between threads is cheap
  - All see the same address space
- Threads become the unit of scheduling
- Processes are now containers in which threads execute

# Processes vs. Threads (2)

---

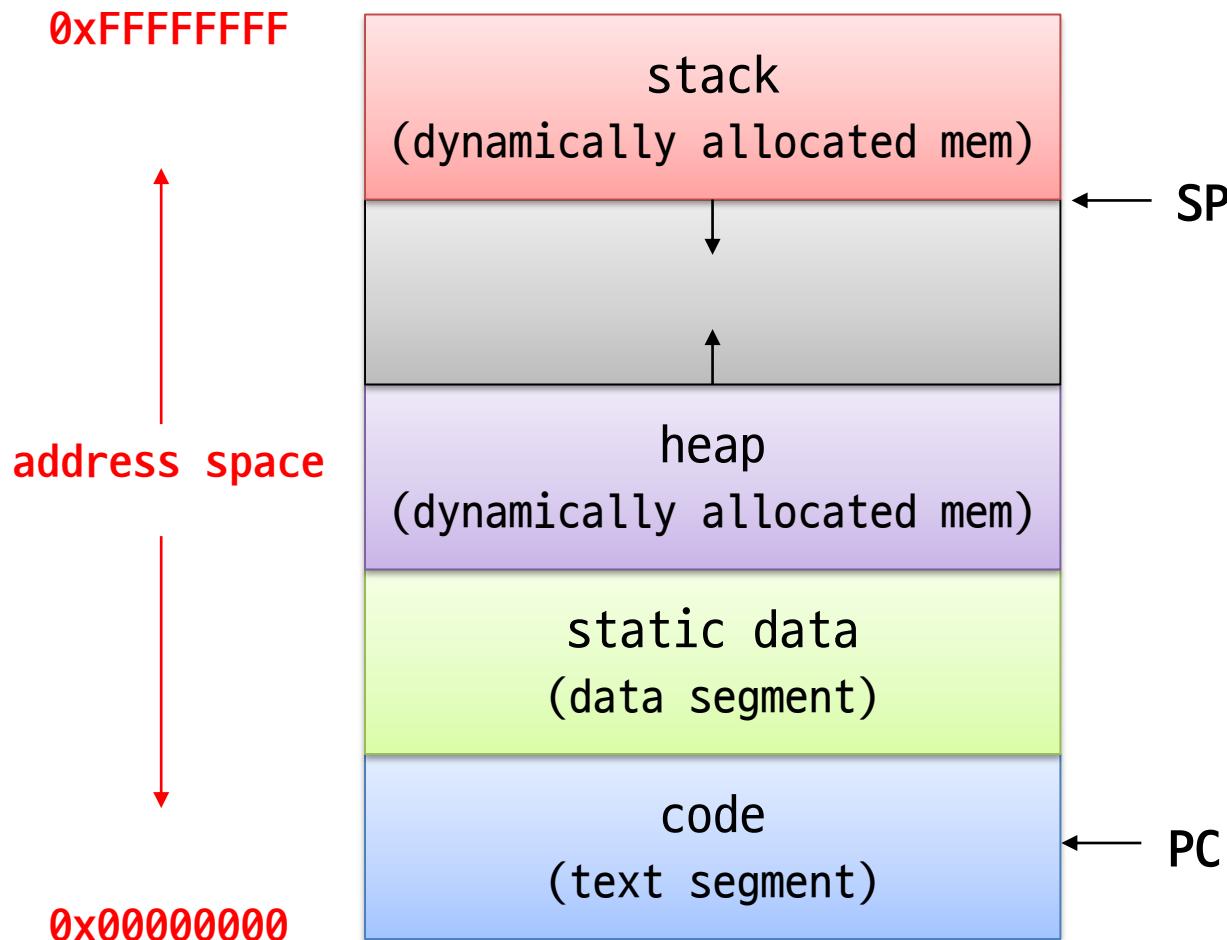
## How threads and processes are similar

- Each has its own logical control flow
- Each can run concurrently with others (possibly on different cores)
- Each is context switched

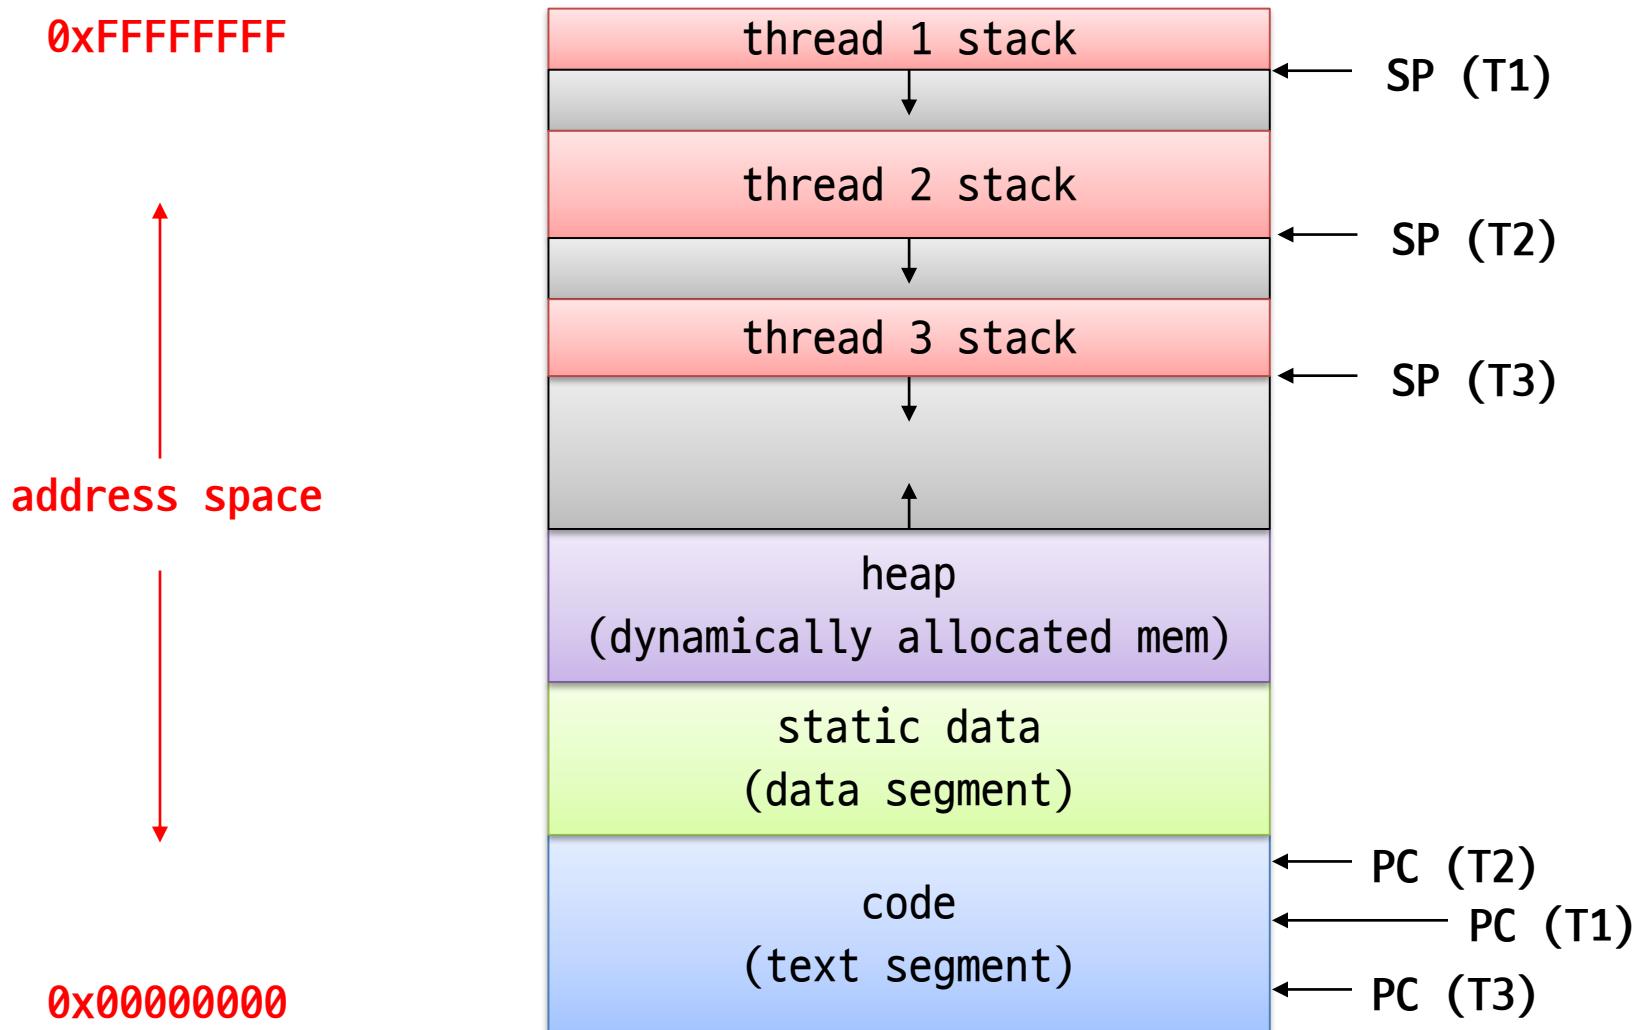
## How threads and processes are different

- Threads share code and some data
  - Processes (typically) do not - use the same code and data copies
- Threads are somewhat less expensive than processes
  - Process control (creating and reaping) is twice as expensive as thread control
  - Linux numbers:
    - ~20K cycles to create and reap a process
    - ~10K cycles (or less) to create and reap a thread

# Process Address Space



# Address Space with Threads



# Classification

# threads per addr space	# of addr spaces	One	Many
One	One	<ul style="list-style-type: none"><li>• MS-DOS</li><li>• Early Macintosh</li></ul>	<ul style="list-style-type: none"><li>• Traditional UNIX</li></ul>
Many	One	<ul style="list-style-type: none"><li>• Many embedded OSes</li><li>• VxWorks</li><li>• uLinux</li></ul>	<ul style="list-style-type: none"><li>• Mach</li><li>• OS/2</li><li>• Linux</li><li>• Windows</li><li>• Mac OS X</li><li>• Solaris</li><li>• HP-UX</li></ul>

# Threads Interface (1)

---

## pthreads

- A **POSIX standard** (IEEE 1003.1c) API for thread creation and synchronization
- API specifies behavior of the thread library
- Implementation is up to development of the library
- Common in UNIX operating systems

# Threads Interface (2)

---

## POSIX-style threads

- pthreads
- DCE threads (early version of pthreads)
- Unix International (UI) threads (Solaris threads)
  - Sun Solaris 2, SCO Unixware 2

## Microsoft-style threads

- Win32 threads
  - Microsoft Windows 98/NT/2000/XP
- OS/2 threads
  - IBM OS/2

# pthreads (1)

## Thread creation/termination

```
int pthread_create (pthread_t *tid,  
                    pthread_attr_t *attr,  
                    void *(start_routine)(void *),  
                    void *arg);
```

```
void pthread_exit (void *retval);
```

```
int pthread_join (pthread_t tid,  
                  void **thread_return);
```

# The Pthreads "hello, world" Program

```
#include <stdio.h>
#include <pthread.h>

void *threadfunc(void *vargp);

/* thread routine */
void *threadfunc(void *vargp) {
    sleep(1);
    printf("Hello, world!\n");
    return NULL;
}

int main() {
    pthread_t tid;

    pthread_create(&tid, NULL, threadfunc, NULL);
    printf("main\n");
    pthread_join(tid, NULL);
    printf("main2\n");
    sleep(2);
    return 0;
}
```

```
# gcc ex.c -lpthread
# ./a.out
main
Hello, world!
main2
```

# pthreads (2)

## Mutexes

```
int pthread_mutex_init  
    (pthread_mutex_t *mutex,  
     const pthread_mutexattr_t *mattr);
```

```
void pthread_mutex_destroy  
    (pthread_mutex_t *mutex);
```

```
void pthread_mutex_lock  
    (pthread_mutex_t *mutex);
```

```
void pthread_mutex_unlock  
    (pthread_mutex_t *mutex);
```

# Threads using shared data

```
#include <pthread.h>
#define MAX_THREAD 20

void *threadcount(void *data) {
    int *count = (int *)data;
    int i;
    for (i=0; i<100; i++) {
        *count = *count+1;
    }
}
int main(int argc, char **argv) {
    pthread_t thread_id[MAX_THREAD];
    int i = 0;
    int count = 0;
    for(i = 0; i < MAX_THREAD; i++) {
        pthread_create(&thread_id[i], NULL, threadcount, (void *)&count);
    }
    for(i = 0; i < MAX_THREAD; i++) {
        pthread_join(thread_id[i], NULL);
    }
    printf("Main Thread : %d\n", count);
    return 0;
}
```

```
# gcc ex.c -lpthread
# ./a.out
Main Thread : 2000
# ./a.out
Main Thread : 1957
```

# pthreads (3)

## Condition variables

```
int pthread_cond_init  
    (pthread_cond_t *cond,  
     const pthread_condattr_t *cattr);
```

```
void pthread_cond_destroy  
    (pthread_cond_t *cond);
```

```
void pthread_cond_wait  
    (pthread_cond_t *cond,  
     pthread_mutex_t *mutex);
```

```
void pthread_cond_signal  
    (pthread_cond_t *cond);
```

```
void pthread_cond_broadcast  
    (pthread_cond_t *cond);
```

# Threading Issues (1)

---

fork() and exec() can be issue

When a thread calls fork()

- Does the new process duplicate all the threads?
- Is the new process single-threaded?

Some UNIX systems support two versions of fork()

- In pthreads,
  - fork() duplicates only a calling thread
- In the Unix international standard,
  - fork() duplicates all parent threads in the child
  - fork1() duplicates only a calling thread

Normally, exec() replaces the entire process

If a thread call exit()?

If the main thread dies(return, exit()) before child threads?

# Threading Issues (2)

---

## Thread cancellation

- The task of terminating a thread before it has completed

### Asynchronous cancellation

- Terminates the target thread immediately
- What happens if the target thread is holding a resource, or it is in the middle of updating shared resources?

### Deferred cancellation

- The target thread is terminated at the cancellation points
- The target thread periodically check if it should be cancelled

pthreads API supports both asynchronous and deferred cancellation

# Threading Issues (3)

---

## Signal handling

- Where should a signal be delivered?

To the thread to which the signal applies

- for synchronous signals

To every thread in the process

To certain threads in the process

- Typically only to a single thread found in a process that is not blocking the signal
- pthreads: per-process pending signals, per-thread blocked signal mask

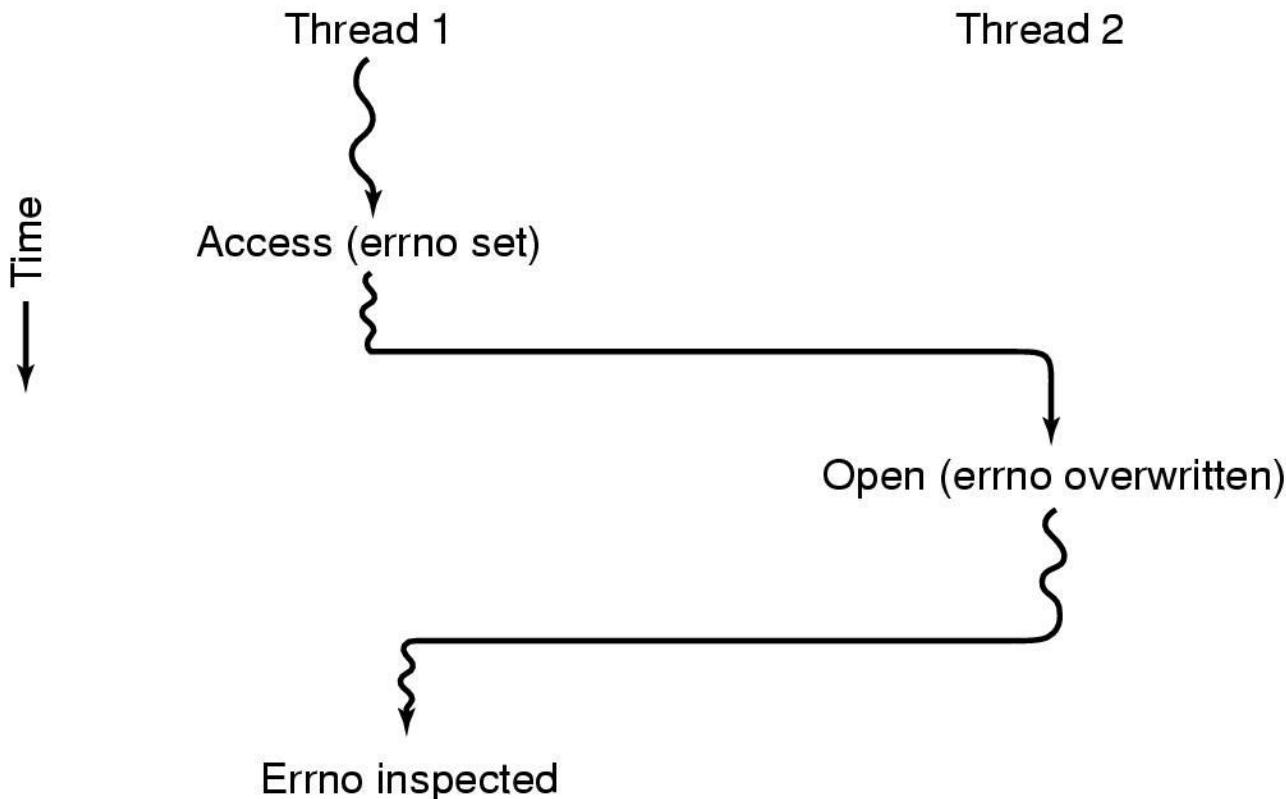
Assign a specific thread to receive all signals for the process

- Solaris 2

# Threading Issues (4)

## Using libraries having internal variables

- errno
  - `#include <errno.h>`
  - Each thread should have **its own independent version** of the `errno` variable



# Threading Issues (4)

## Multithread-safe (MT-safe)

- A set of functions can be said to be multithread-safe or reentrant, when the functions may be called by more than one thread at a time
- Functions that access no global data or read-only global data are trivially MT-safe
- Functions that modify global state must be made MT-safe by synchronizing access to the shared data

Threads can have private global variables

