

Threads Implementation

Jo, Heeseung

Today's Topics

How to implement threads?

- User-level threads
- Kernel-level threads

Which one is the fastest?

Get factorial value of N

- iterative factorial
- recursive factorial

iterative

local function

- myswap()

가

library function

- strcpy()

system call

- getpid()

Kernel/User-level Threads

Who is responsible for creating/managing threads?

- The OS ([kernel-level threads](#))
 - Thread creation and management requires system calls
- The user-level process ([user-level threads](#))
 - A library linked into the program manages the threads

Why is user-level thread management possible?

- Threads share the same address space
 - The thread manager doesn't need to manipulate address spaces
- Threads only differ in hardware contexts (roughly)
 - PC, SP, registers
 - These can be manipulated by the user-level process itself

Kernel-level Threads (1)

OS-managed threads

- The OS manages threads and processes
- All thread operations are implemented in the kernel
- The OS schedules all of the threads in a system
 - If one thread in a process blocks (e.g., on I/O), the OS knows about it, and can run other threads from that process
 - Possible to overlap I/O and computation inside a process
- Kernel threads are cheaper than processes
 - Less state to allocate and initialize
- Windows, Solaris, Tru64 Unix, Linux, MacOS

Kernel-level Threads (2)


Limitations

- They can still be **too expensive**
 - For fine-grained concurrency, we need even cheaper threads
 - Ideally, we want thread operations as fast as a procedure call
- **Thread operations are all system calls**
 - The program must cross an extra protection boundary on every thread operation
 - Even when the processor is being switched between threads in the same address space
 - The OS must perform all of the usual argument checks
- **Must maintain kernel state for each thread**
 - Can place limit on the number of simultaneous threads
 - In Linux, 256430 (`/proc/sys/kernel/threads-max`)
- Kernel-level threads have to be general to support the needs of all programmers, languages, runtime systems, etc.

Implementing Kernel-level Threads

Kernel-level threads

- Kernel-level threads are similar to original process management and implementation



User-level Threads (1)

Motivation



- To make threads **cheap and fast**, they need to be implemented at the user level
- **Portable**: User-level threads are managed entirely by the runtime system (user-level library)

User-level threads are **small and fast**.

- Each thread is represented simply by a PC, registers, a stack, and a **small thread control block (TCB)**
- Creating a thread, switching between threads, and synchronizing threads are done **via procedure calls** (No kernel involvement)
- User-level thread operations can be **10-100x faster** than kernel-level threads

Implementing User-level Threads (1)

User-level threads




Kernel thread

Implementing User-level Threads (2)

Thread context switch


- Very simple for user-level threads
- Save context of currently running thread
 - : **push all machine state onto its stack**
- Restore context of the next thread
 - : **pop machine state from next thread's stack**
- The next thread becomes the current thread
- Return to caller as the new thread
 - : execution resumes at PC of next thread



User-level Threads (2)

Limitations

- User-level threads are **invisible to the OS**
 - They are not well integrated with the OS
- As a result, **the OS can make poor decisions**
 - Scheduling a process with only idle threads
 - Blocking a process whose thread initiated I/O, even though the process has other threads that are ready to run
 - Unscheduling a process with a thread holding a lock
- Solving this requires coordination between the kernel and the user-level thread manager
 - e.g., all blocking system calls should be emulated in the library via non-blocking calls to the kernel



Implementing User-level Threads (3)

Thread scheduling

- A thread scheduler determines when a thread runs
 - Just like the OS and processes
 - But implemented at user-level as a library
- Queues to keep track of what threads are doing
 - Run queue: threads currently running
 - Ready queue: threads ready to run
 - Wait queue: threads blocked for some reason (maybe blocked on I/O or a lock)
- How can we prevent a thread from hogging the CPU?

Implementing User-level Threads (4)

Non-preemptive scheduling

- Force everybody to cooperate
 - Threads willingly give up the CPU by `calling yield()`
- `yield()` calls into the scheduler, which context switches to another ready thread

```
Thread ping ()  
{  
    while (1) {  
        printf ("ping\n");  
        yield();  
    }  
}
```

```
Thread pong ()  
{  
    while (1) {  
        printf ("pong\n");  
        yield();  
    }  
}
```

- What happens if a thread never calls `yield()`?


Implementing User-level Threads (5)

Preemptive scheduling

- Need to regain control of processor asynchronously
- Scheduler requests that a timer interrupt be delivered by the OS periodically
 - Usually delivered as a UNIX signal
 - Signals are just like software interrupts, but delivered to user-level by the OS instead of delivered to OS by hardware
- At each timer interrupt, scheduler gains control and context switches as appropriate

Case study : Go Lang

3rd Language to learn (2020)



Go is developed for Google infra structure

- Before Go, C++ is used
- C++ is fast, but every day code update, complex, long long build time

Case study : Go Lang

Pros

- Low memory consumption : 2KB stack
- Cheap creation and deletion of threads
 - goroutine : user-level thread (green thread)
- Cheap context switch cost


Go runtime scheduler

- Started with a Go program
- scheduling goroutine

Scheduler rule

- Kernel thread is cost, and so minimize them
- High concurrency with many many goroutine
- N goroutine on N cores

Case study : Go Lang



Case study : Go Lang

Scheduler concepts

- Reuse threads
 - Unused go thread -> idle state
 - After sometime in idle, remove
- Limit threads accessing runqueue
 - Max number of threads is 10000 (default, modifiable)
 - Accessing local and global run queue needs lock
- Distributed runqueues
 - Minimize schedule lock